

Data structures

- Organize your data to support various queries using little time and/or space

- Given n elements $A[1..n]$
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$
- Trivial solution: scan A . Takes time $\Theta(n)$
- Best possible given A, x .
- What if we are first given A , are allowed to **preprocess** it, can we then answer SEARCH queries faster?
- How would you preprocess A ?

- Given n elements $A[1..n]$
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$
- Preprocess step: Sort A . Takes time $O(n \log n)$, Space $O(n)$
- $\text{SEARCH}(A[1..n], x) :=$ /* Binary search */
 If $n = 1$ then return YES if $A[1] = x$, and NO otherwise
 else
 if $A[n/2] \leq x$ then return $\text{SEARCH}(A[n/2..n])$
 else return $\text{SEARCH}(A[1..n/2])$
- Time $T(n) = ?$

- Given n elements $A[1..n]$
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$
- Preprocess step: Sort A . Takes time $O(n \log n)$, Space $O(n)$
- $\text{SEARCH}(A[1..n], x) :=$ /* Binary search */
 If $n = 1$ then return YES if $A[1] = x$, and NO otherwise
 else
 if $A[n/2] \leq x$ then return $\text{SEARCH}(A[n/2..n])$
 else return $\text{SEARCH}(A[1..n/2])$
- Time $T(n) = O(\log n)$.

- Given n elements $A[1..n]$ each $\leq k$, can you do faster?
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A$?
- DIRECT ADDRESS:
- Preprocess step: Initialize $S[1..k]$ to 0
 For $(i = 1 \text{ to } n) S[A[i]] = 1$
- $T(n) = O(n)$, Space $O(k)$
- $\text{SEARCH}(A, x) = ?$

- Given n elements $A[1..n]$ each $\leq k$, can you do faster?
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$
- DIRECT ADDRESS:
- Preprocess step: Initialize $S[1..k]$ to 0
 For $(i = 1 \text{ to } n) S[A[i]] = 1$
- $T(n) = O(n)$, Space $O(k)$
- $\text{SEARCH}(A, x) = \text{return } S[x]$
- $T(n) = O(1)$

- Dynamic problems:
- Want to support SEARCH, INSERT, DELETE
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$

- If numbers are small, $\leq k$

Preprocess: Initialize S to 0.

$\text{SEARCH}(x) := \text{return } S[x]$

$\text{INSERT}(x) := \dots??$

$\text{DELETE}(x) := \dots??$

- Dynamic problems:
- Want to support SEARCH, INSERT, DELETE
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$

- If numbers are small, $\leq k$

Preprocess: Initialize S to 0.

$\text{SEARCH}(x) := \text{return } S[x]$

$\text{INSERT}(x) := S[x] = 1$

$\text{DELETE}(x) := S[x] = 0$

- Time $T(n) = O(1)$ per operation
- Space $O(k)$

- Dynamic problems:
- Want to support SEARCH, INSERT, DELETE
- Support $\text{SEARCH}(A, x) := \text{is } x \text{ in } A?$
- What if numbers are not small?
- There exist a number of data structure that support each operation in $O(\log n)$ time
- Trees: AVL, 2-3, 2-3-4, B-trees, red-black, AA, ...
- Skip lists, deterministic skip lists,
- Let's see binary search trees first

Binary tree

Vertices, aka nodes = {a, b, c, d, e, f, g, h, i}

Root = a

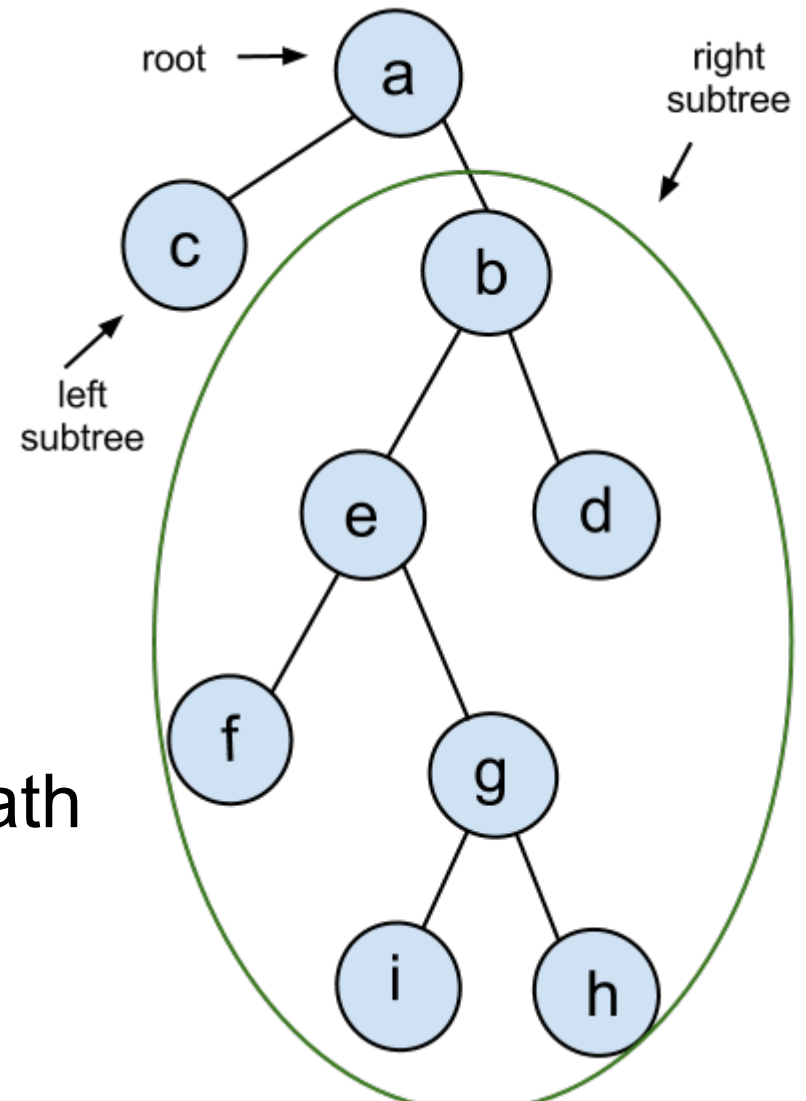
Left subtree = {c}

Right subtree = {b, d, e, f, g, h, i}

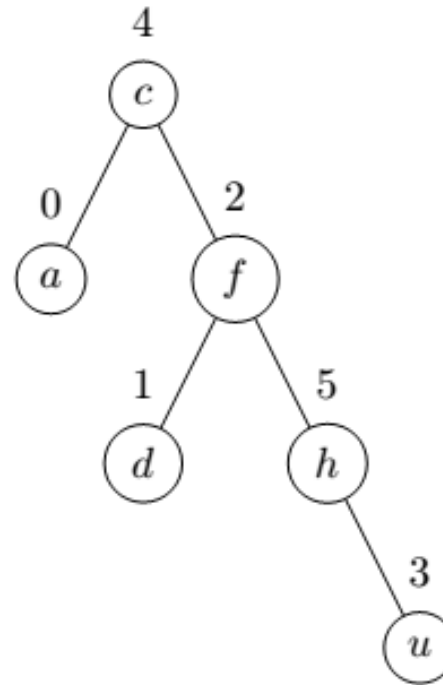
Parent(b) = a

Leaves = nodes with no children
= {c, f, i, h, d}

Depth = length of longest root-leaf path
= 4



How to represent a binary tree using arrays



Index	0	1	2	3	4	5
Key	a	d	f	u	c	h
Parent	4	2	4	5	NULL	2
LeftChild	NULL	NULL	1	NULL	0	NULL
RightChild	NULL	NULL	5	NULL	2	3

Root = 4

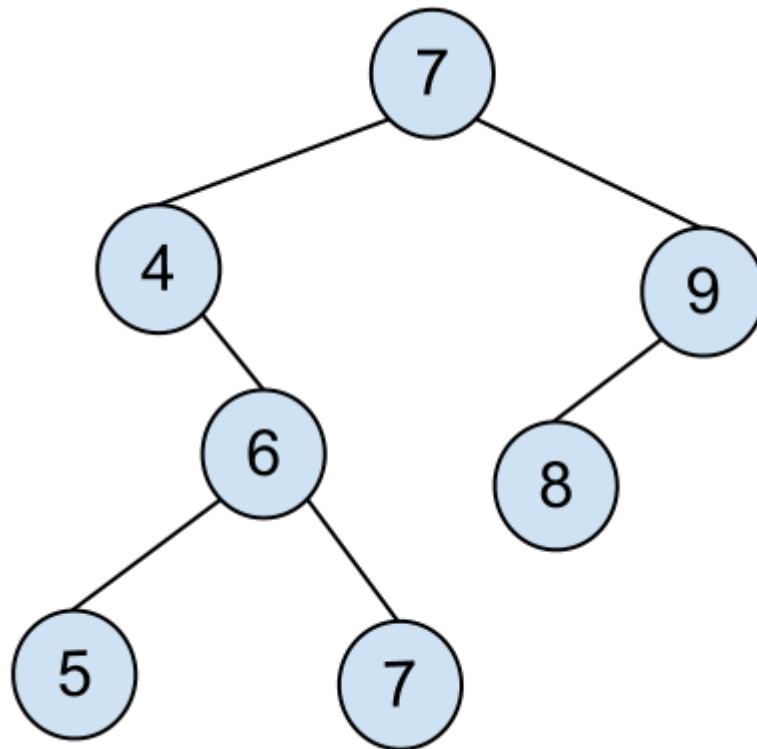
NumNodes = 5

Binary Search Tree is a data structure where we store data in nodes of a binary tree and refer to them as **key** of that node.

The keys in a binary search tree satisfy the **binary search tree property**:

Let $x, y \in V$, if y is in left subtree of $x \implies \text{key}(y) \leq \text{key}(x)$
if y is in right subtree of $x \implies \text{key}(x) < \text{key}(y)$.

Example:



Tree-search(x,k) \ \ Looks for k in tree rooted at x

if x = NULL or k = Key[x]

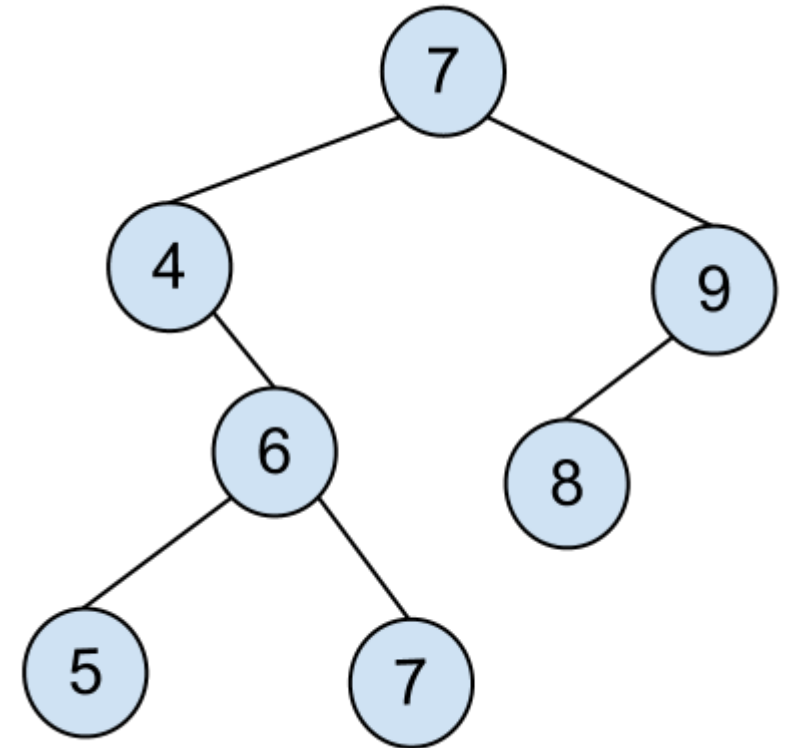
return x

if $k \leq \text{key}[x]$

return Tree-search(LeftChild[x],k)

else

return tree-search(RightChild[x],k)



Running time = $O(\text{Depth})$

Depth = $O(\log n) \Rightarrow$ search time $O(\log n)$

Tree-Search is a generalization of binary search in an array that we saw before.

A sorted array can be thought of as a balanced tree (we'll return to this)

Trees make it easier to think about **inserting** and removing

Insert(k) // Inserts k

 If the tree is empty

 Create a root with key k and return

 Let y be the last node visited during Tree-Search(Root,k)

 If $k \leq \text{Key}[y]$

 Insert new node with key k as left child of y

 If $k > \text{Key}[y]$

 Insert new node with key k as right child of y

Running time = $O(\text{Depth})$

Depth = $O(\log n) \Rightarrow$ insert time $O(\log n)$

Let us see the code in more detail

```

Insert(k):
//If there is no room, do nothing
if NumNodes >= MAXNODES
    return
//Otherwise puts a new node at the end of the arrays
Key[NumNodes] ← k
LeftChild[NumNodes] ← NULL
RightChild[NumNodes] ← NULL
//It remains to determine the parent
//If tree is empty then there is none
if NumNodes = 0
    Root ← 0
    Parent[NumNodes] ← NULL
    NumNodes++
    return
//Otherwise looks for the parent in the tree
x ← Root
forever
    //two ifs to check if x is parent, otherwise search
    if  $k \leq \text{Key}[x]$  and LeftChild[x] = NULL
        LeftChild[x] ← NumNodes
        Parent[NumNodes] ← x
        NumNodes++
        return
    if  $k > \text{Key}[x]$  and RightChild[x] = NULL
        RightChild[x] ← NumNodes
        Parent[NumNodes] ← x
        NumNodes++
        return
    if  $k \leq \text{Key}[x]$  and LeftChild[x]  $\neq$  NULL
        x ← LeftChild[x]
    if  $k > \text{Key}[x]$  and RightChild[x]  $\neq$  NULL
        x ← RightChild[x]

```


Recall want search, insert, and delete in time $O(\log n)$

We need to keep the depth to $O(\log n)$

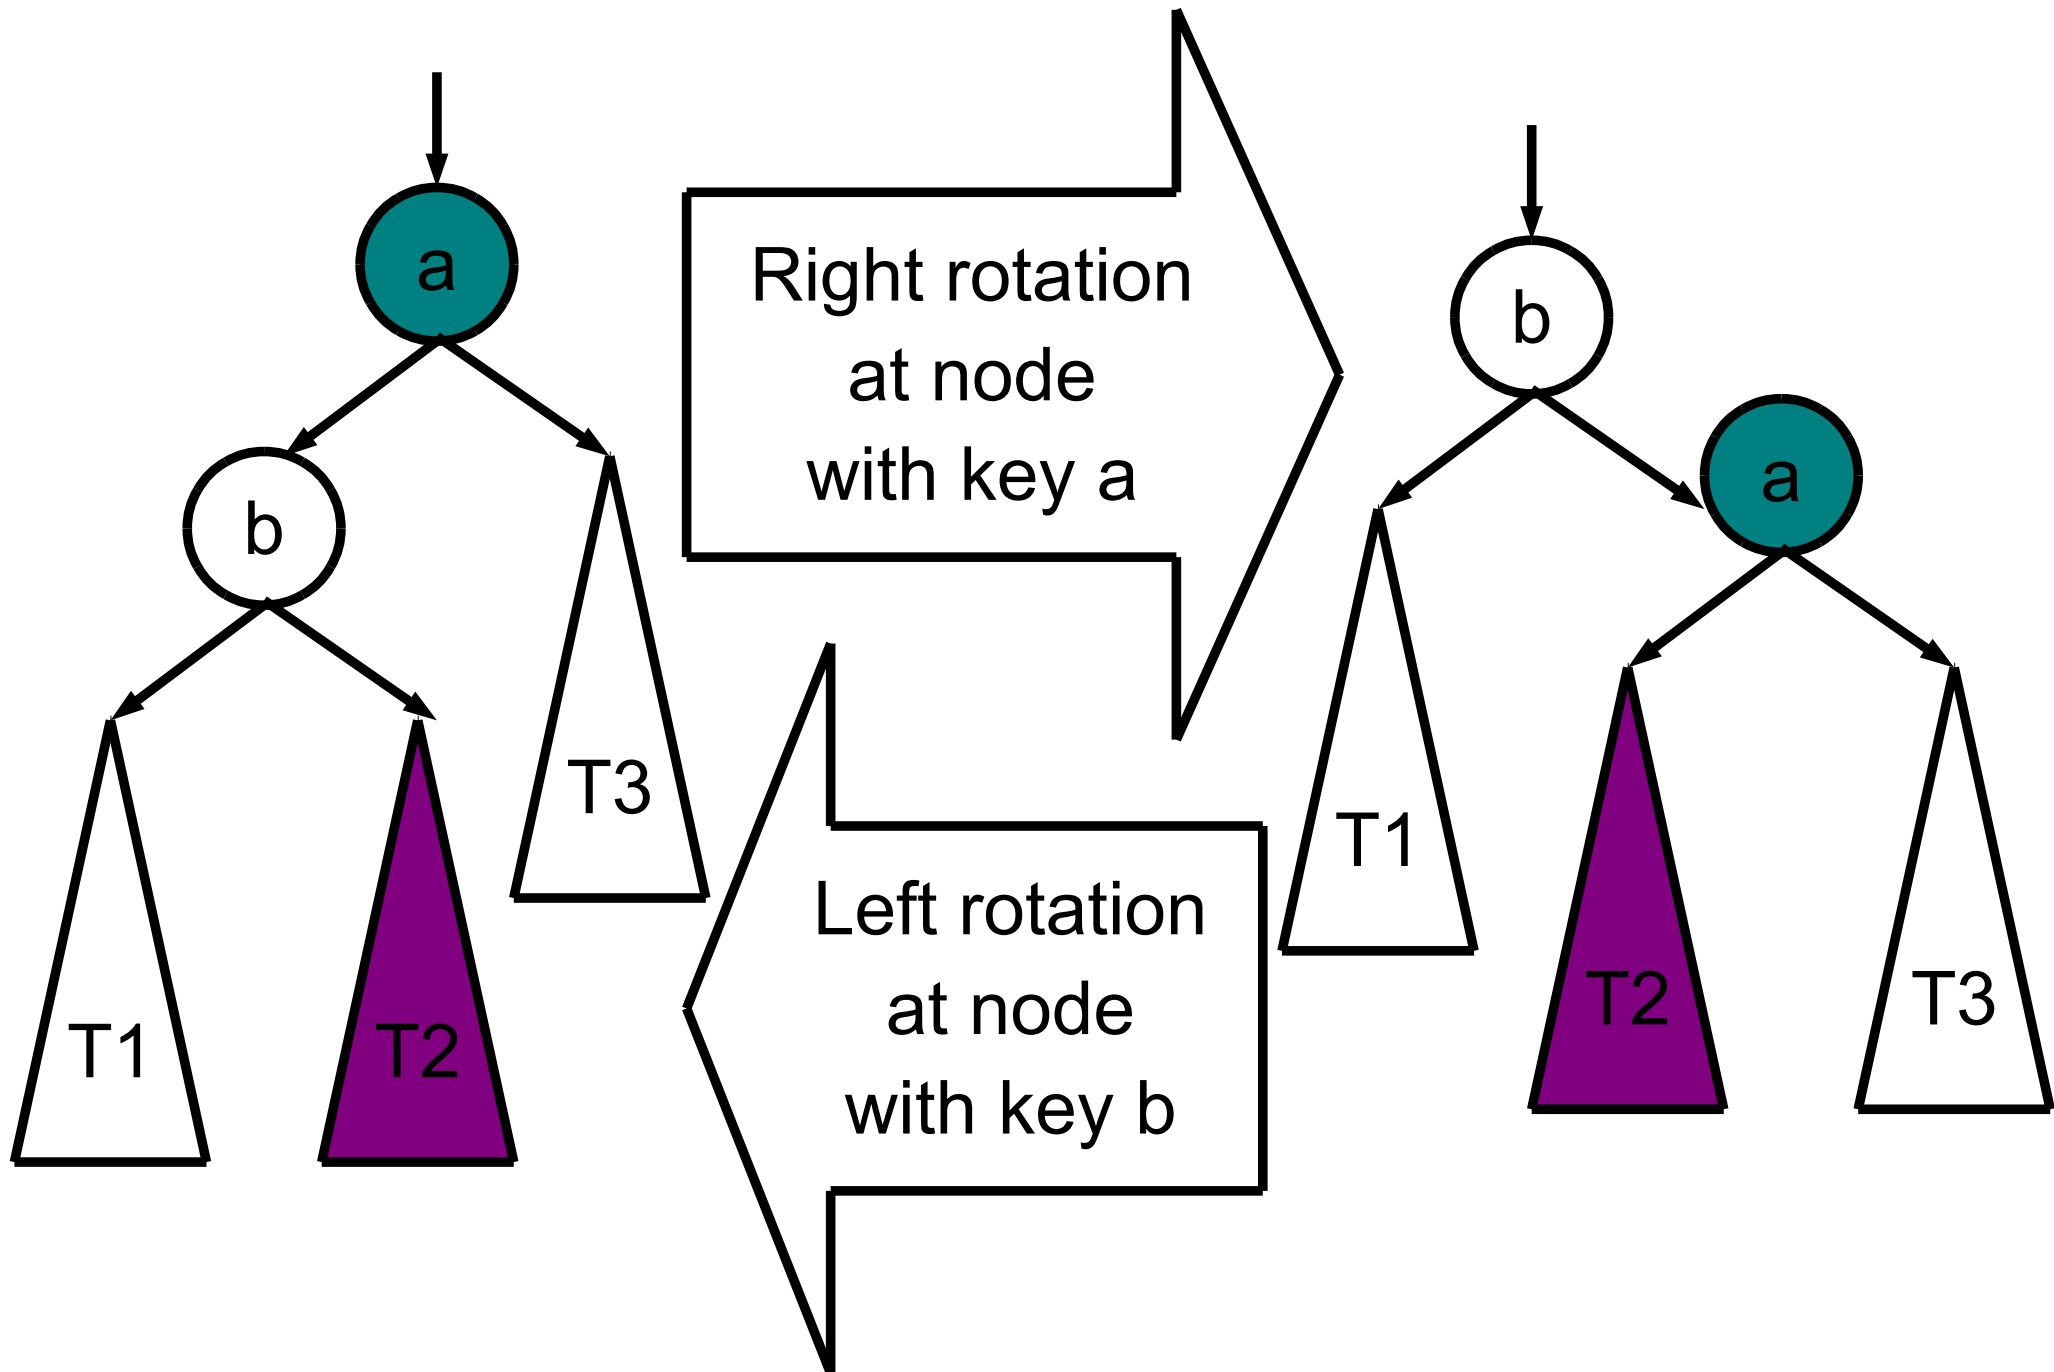
When inserting and deleting, the depth may change.

Must restructure the tree to keep depth $O(\log n)$

A basic restructuring operation is a **rotation**

Rotation is then used by more complicated operations

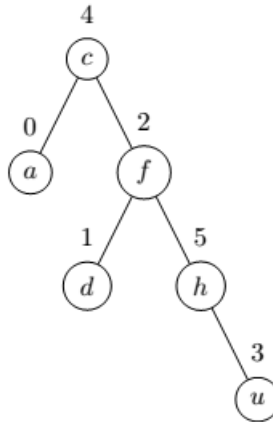
Tree rotations



Tree rotations, code using our representations

```
Rotate-Right(i):  
    if i does not have a left child  
        return  
    L ← LeftChild[i]  
    if i is the root  
        Root ← L, Parent[L] ← NULL  
    If i is a right child of P  
        RightChild[P] ← L, Parent[L] ← P  
    If i is a left child of P  
        LeftChild[P] ← L, Parent[L] ← P  
    LR ← RightChild[L]  
    RightChild[L] ← i  
    Parent[i] ← L  
    LeftChild[i] ← LR  
    If LR ≠ NULL  
        Parent[LR] ← i
```

Tree rotations, code using our representations

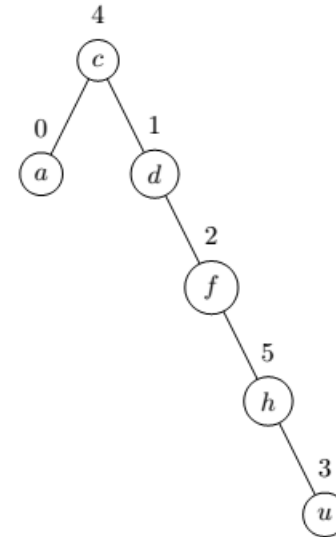


Index	0	1	2	3	4	5
Key	a	d	f	u	c	h
Parent	4	2	4	5	NULL	2
LeftChild	NULL	NULL	1	NULL	0	NULL
RightChild	NULL	NULL	5	NULL	2	3

Root = 4

NumNodes = 5

Figure 4.1: A binary tree and its representation



Index	0	1	2	3	4	5
Key	a	d	f	u	c	h
Parent	4	4	1	5	NULL	2
LeftChild	NULL	NULL	NULL	NULL	0	NULL
RightChild	NULL	2	5	NULL	1	3

Root = 4

NumNodes = 5

Figure 4.2: The effect of RotateRight(2) on the tree in Figure 4.1.

Using rotations to keep the depth
small

- **AVL trees**: binary trees. In any node, heights of children differ by ≤ 1 . Maintain by rotations
- **2-3-4 trees**: nodes have 1, 2, or 3 keys and 2, 3, or 4 children. All leaves same level. To insert in a leaf: add a child. If already 4 children, split the node into one with 2 children and one with 4, add a child to the parent recursively. When splitting the root, create new root.

Deletion is more complicated.

- **B-trees**: a generalization of 2-3-4 trees where can have more children. Useful in some disk applications where loading a node corresponds to reading a chunk from disk
- **Red-black trees**: A way to “simulate” 2-3-4 trees by a binary tree. E.g. split 2 keys in same 2-3-4 node into 2 red-black nodes. Color edges red or black depending on whether the child comes from this splitting or not, i.e., is a child in the 2-3-4 tree or not.

We see in detail what may be the simplest variant of these:

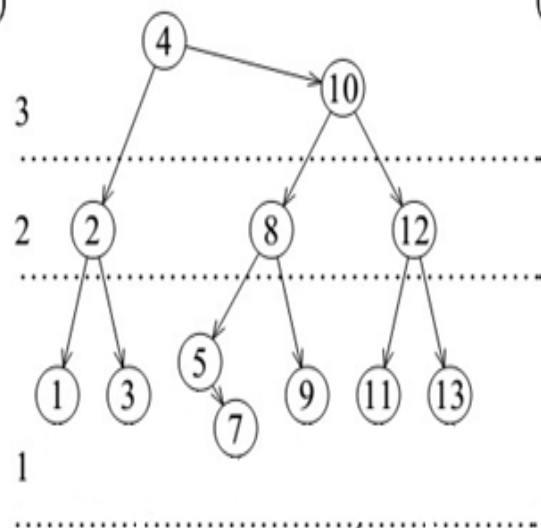
AA Trees

First we see pictures,

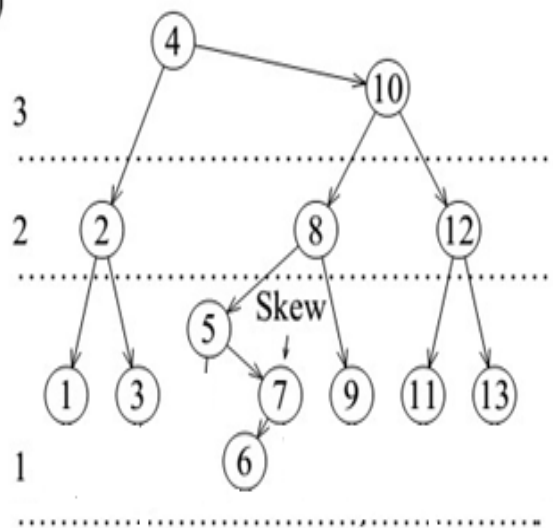
then formalize it,

then go back to pictures.

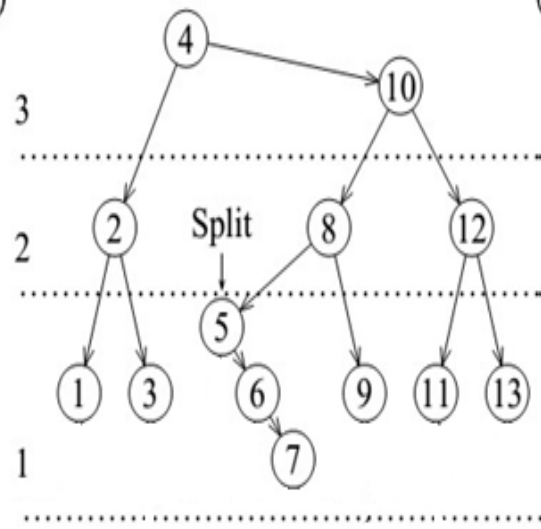
(a)



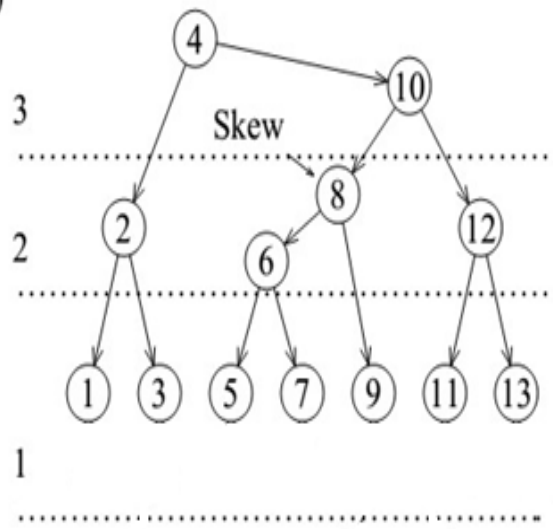
(b)



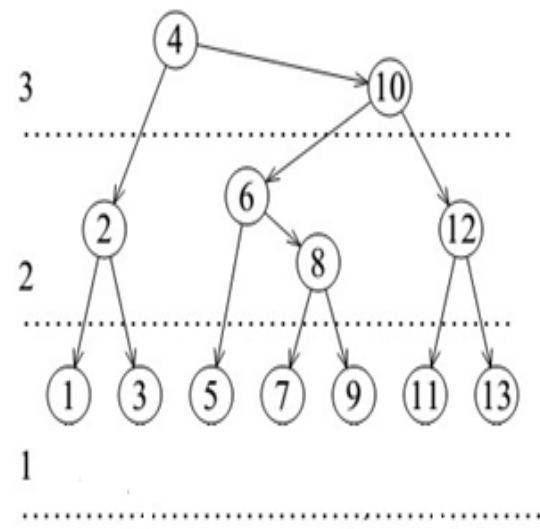
(c)



(d)



(e)



- **Definition:** An **AA Tree** is a binary search tree where each node has a **level**, satisfying:

- (1) The level of every leaf node is one.

- (2) The level of every left child is exactly one less than that of its parent.

- (3) The level of every right child is equal to or one less than that of its parent.

- (4) The level of every right grandchild is strictly less than that of its grandparent.

- (5) Every node of level greater than one has two children.

- Intuition: “the only path with nodes of the same level is a single left-right edge”

- **Fact:** An **AA Tree** with n nodes has depth $O(\log n)$

- **Proof:**

Suppose the tree has depth d .

The **level** of the root is at least $d/2$.

Since every node of level > 1 has two children, the tree contains a full binary tree of depth at least $d/2 - 1$. Such a tree has at least $2^{d/2 - 1}$ nodes. ■

- Restructuring an AA tree after an addition:
- Rule of thumb:

First make sure that only left-right edges are within nodes of the same level (Skew)

then worry about length of paths within same level (Split)

Restructuring operations:

Skew(x): If x has left-child with same level

RotateRight(x)

Split(x): If the level of the right child of the right child of x
is the same as the level of x,

Level[RightChild[x]]++;

RotateLeft(x)

AA-Insert(k):

Insert k as in a binary search tree

/* For every node from new one back to root,
do skew and split

*/

$x \leftarrow \text{NumNodes}-1$ //New node is last in array

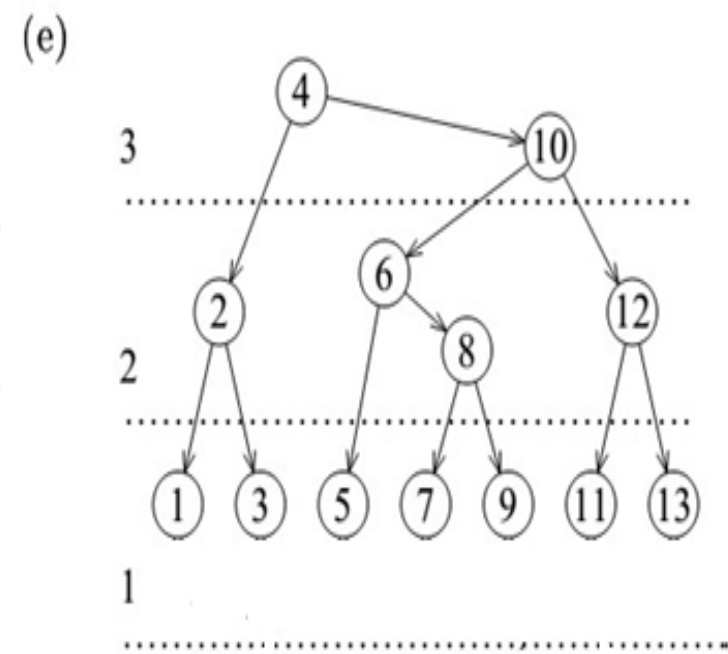
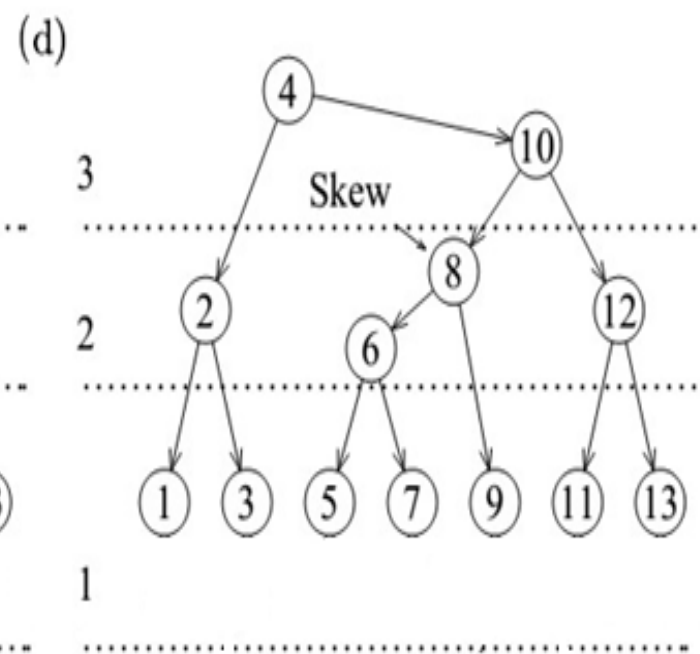
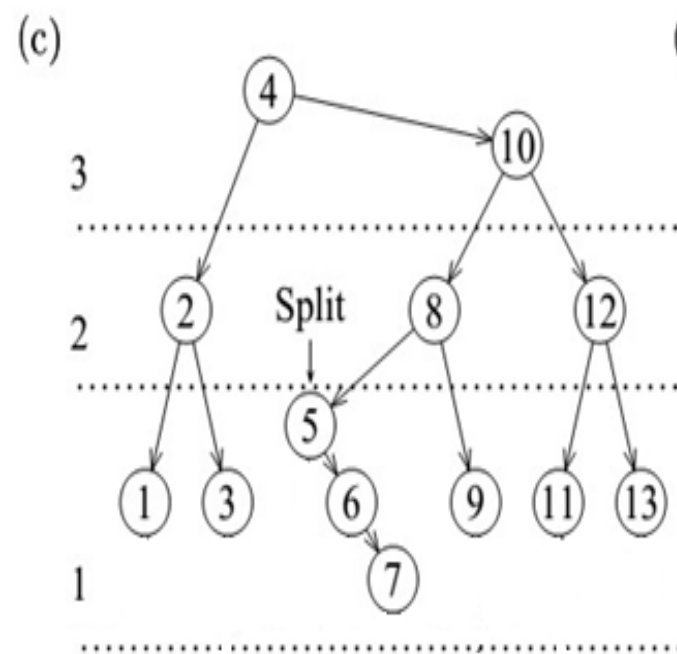
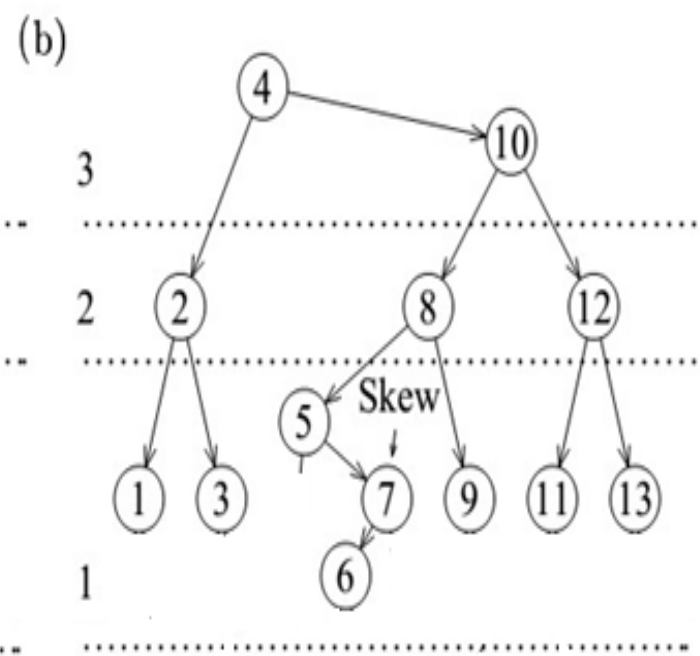
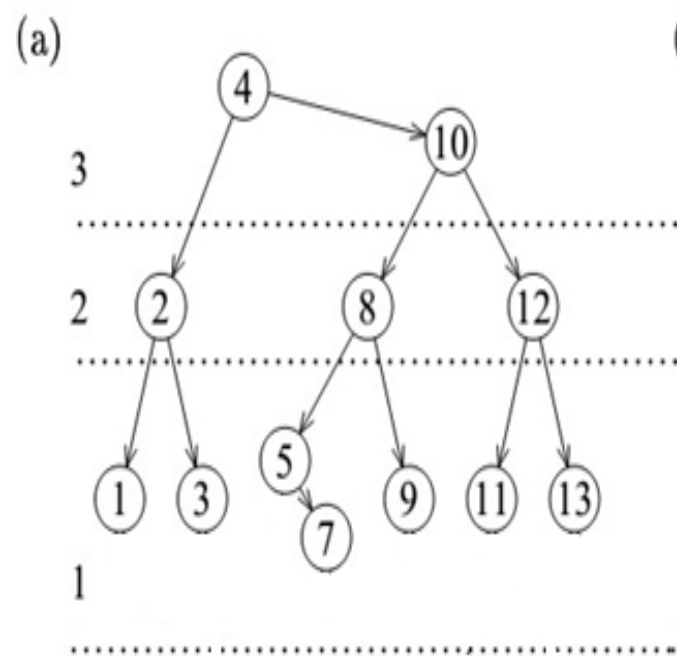
while $x \neq \text{NULL}$

Skew(x)

Split(x)

$x \leftarrow \text{Parent}[x]$

Inserting 6



Deleting in an AA tree:

Decrease Level(x):

If one of x's children is two levels below x,
decrease the level of x by one.

If the right child of x had the same level of x, decrease the
level of the right child of x by one too.

Delete(x): Suppose x is a leaf

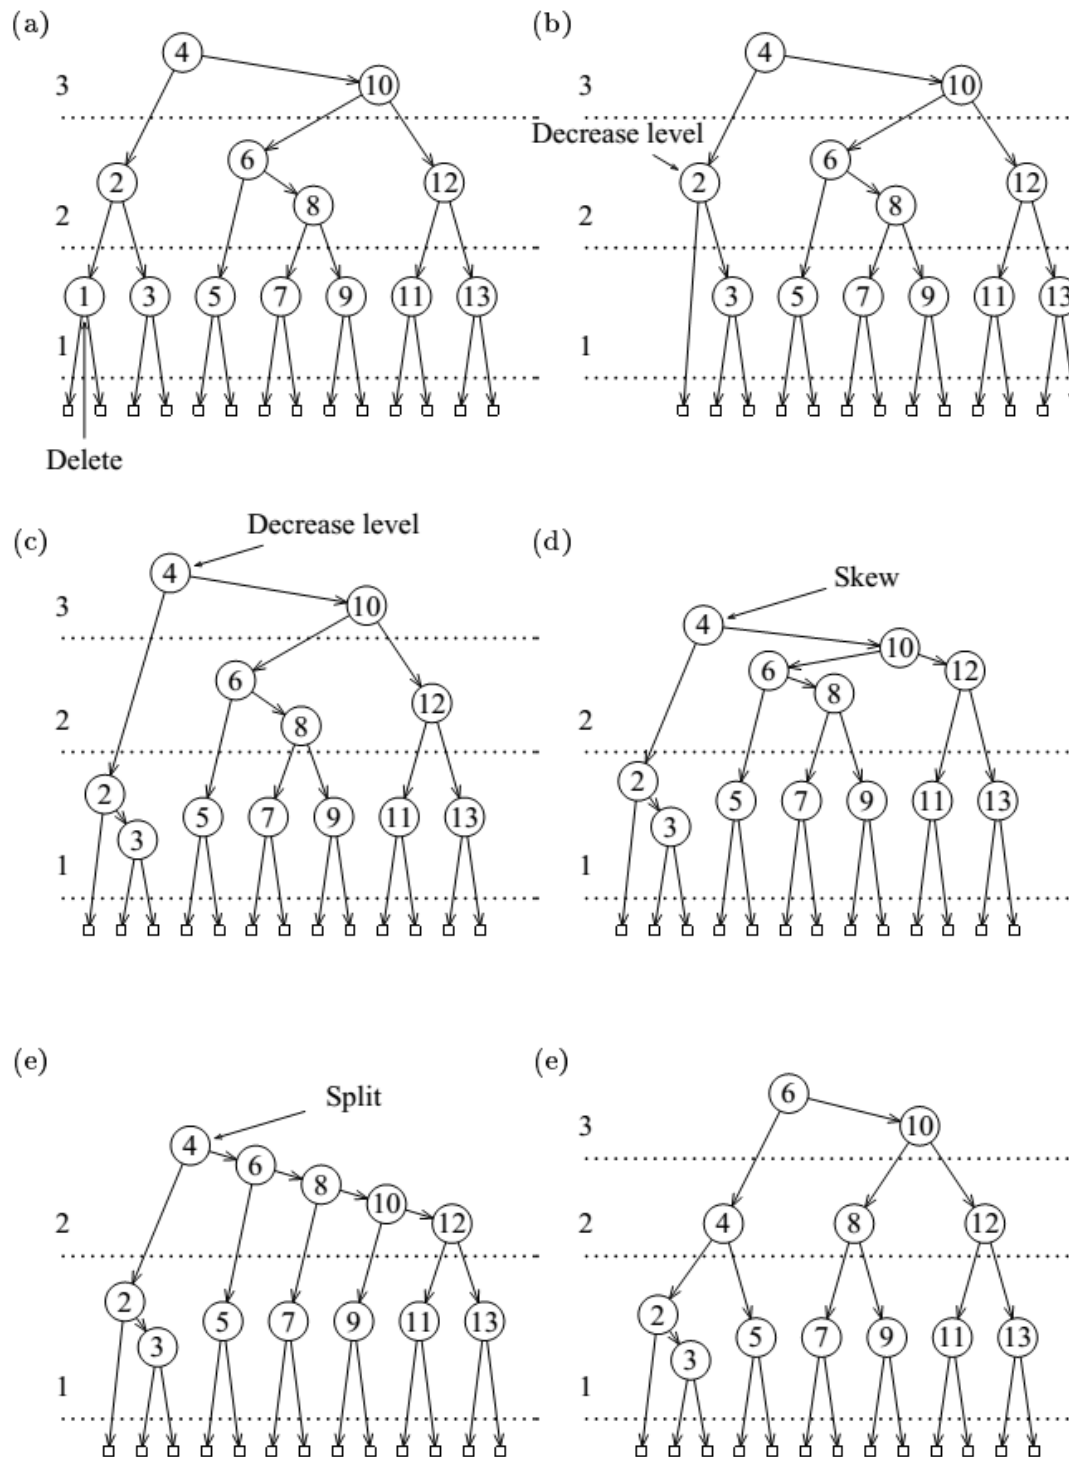
Delete x.

Follow the path from x to the root and at each node y do:

Decrease level(y).

Skew(y); Skew(y.right); Skew(y.right.right);

Split(y); Split(y.right);



Rotate right 10,
get $8 \leftarrow 10$,
so again
rotate right 10

Fig.2. Example of deletion.

Delete(x):

If x is not a leaf, find the smallest leaf bigger than x.key, swap it with x, and remove that leaf.

To find that leaf, just perform search, and when you hit x go, e.g., right. It's the same thing as searching for $x.key + \epsilon$

So swapping these two won't destroy the tree properties

Remark about memory implementation:

Could use new/malloc free/dispose to add/remove nodes.

However, this may cause memory segmentation.

It is possible to implement any tree using array A so that:
at any point in time, if n elements are in the tree, those will take elements $A[1..n]$ in the array only.

To do this, when you remove node with index i in the array, swap $A[i]$ and $A[n]$. Use parent's pointers to update.

Summary

Can support SEARCH, INSERT, DELETE in time $O(\log n)$ for arbitrary keys

Space: $O(n)$. For each key we need to store level and pointers to children, and possibly pointer to parent too.

Can we achieve space closer to n , like $n + 0.001n$?

Surprisingly, this is possible:

Optimal Worst-Case Operations for Implicit Cache-Oblivious Search Trees, by Franceschini and Grossi

Hash functions

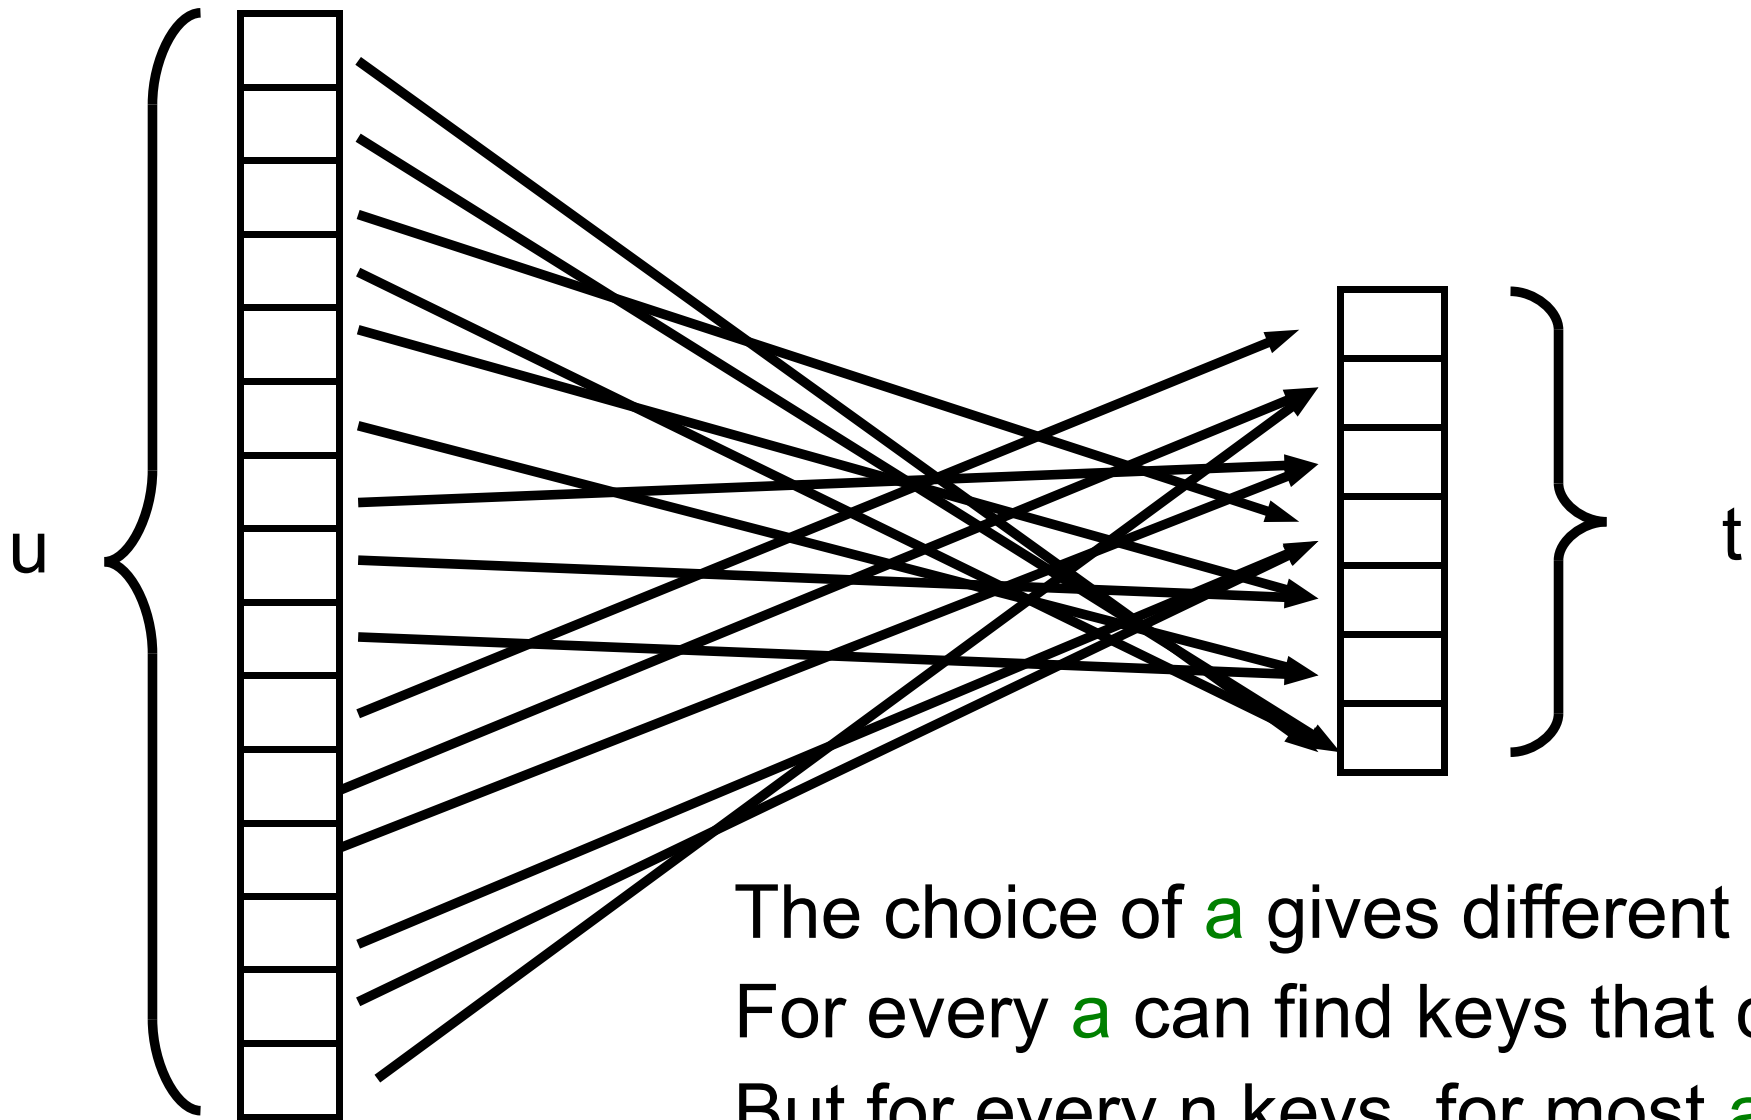
- We have seen how to support SEARCH, INSERT, and DELETE in time $O(\log n)$ and space $O(n)$ for arbitrary keys
- If the keys are small integers, say in $\{1, 2, \dots, t\}$ for a small t we can do it in time ?? and space ??

- We have seen how to support SEARCH, INSERT, and DELETE in time $O(\log n)$ and space $O(n)$ for arbitrary keys
- If the keys are small integers, say in $\{1, 2, \dots, t\}$ for a small t we can do it in time $O(1)$ and space $O(t)$
- Can we have the same time for arbitrary keys?
- Idea: Let's make the keys small.

Keys

Hash function h_a

Table

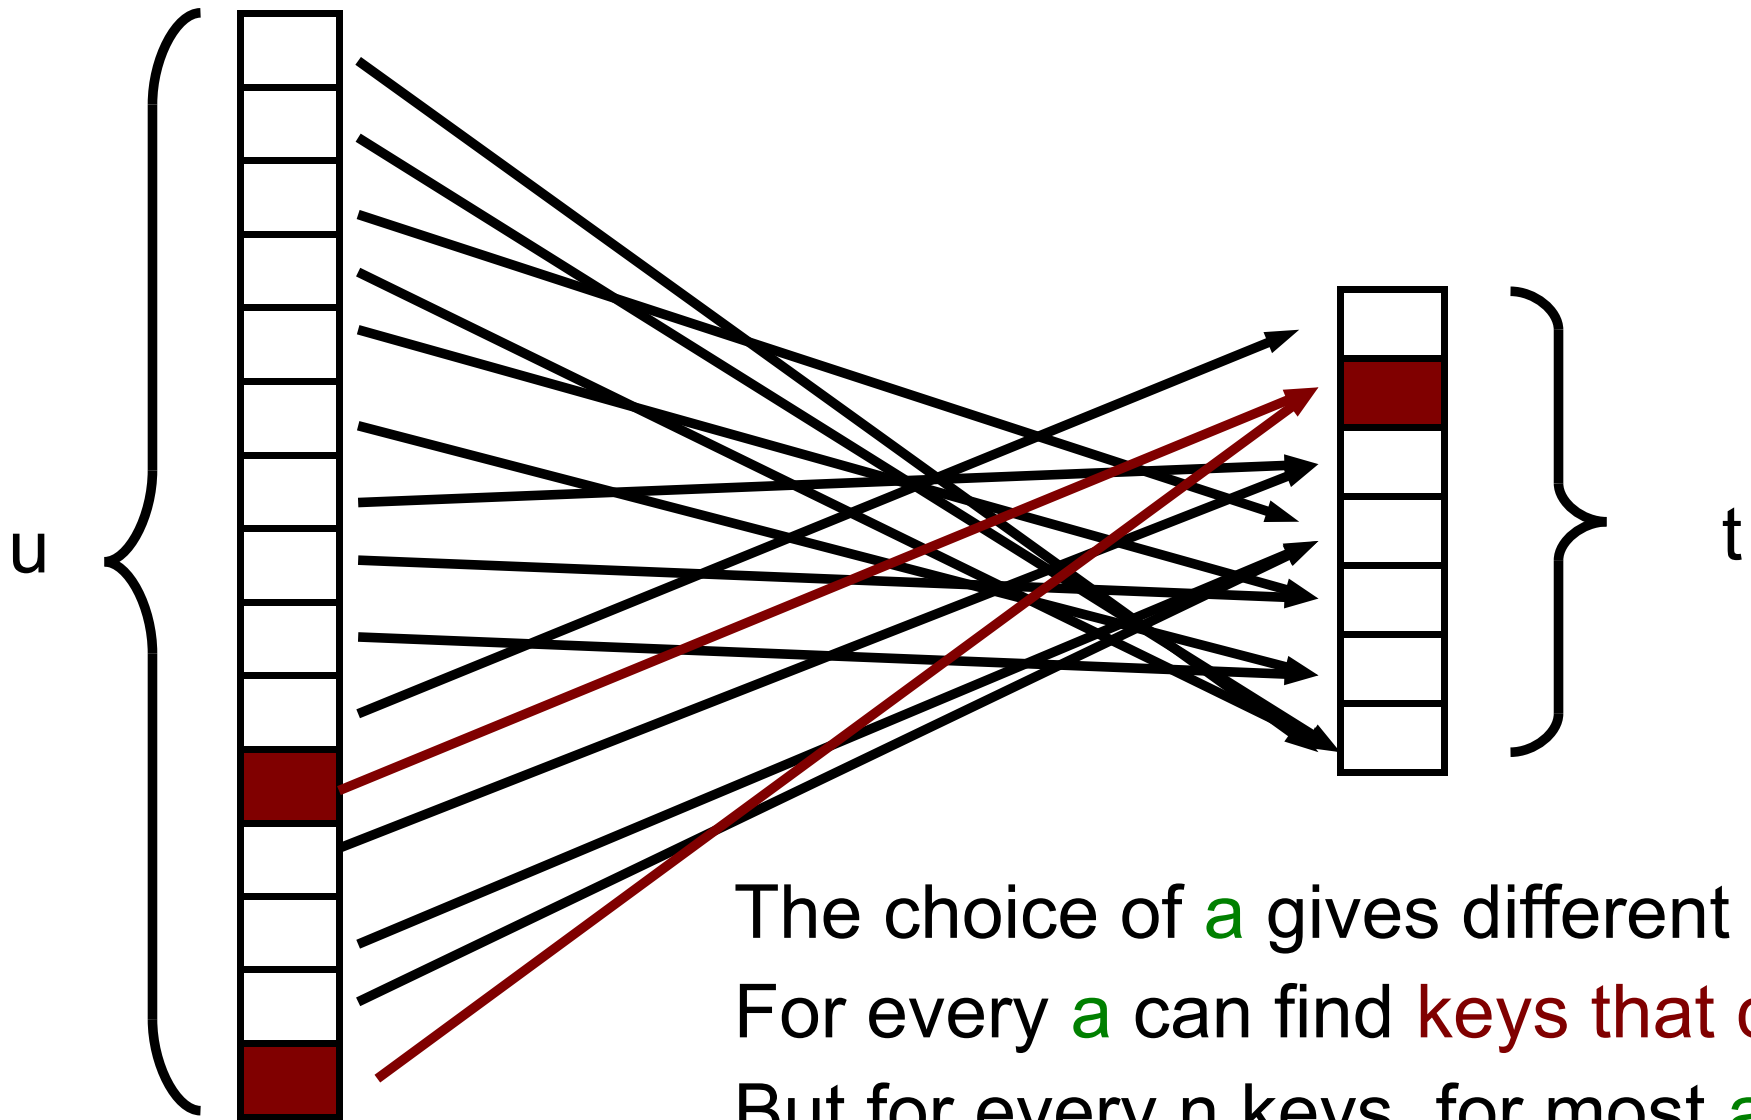


The choice of a gives different arrows
For every a can find keys that collide
But for every n keys, for most a
there are no collisions

Keys

Hash function h_a

Table

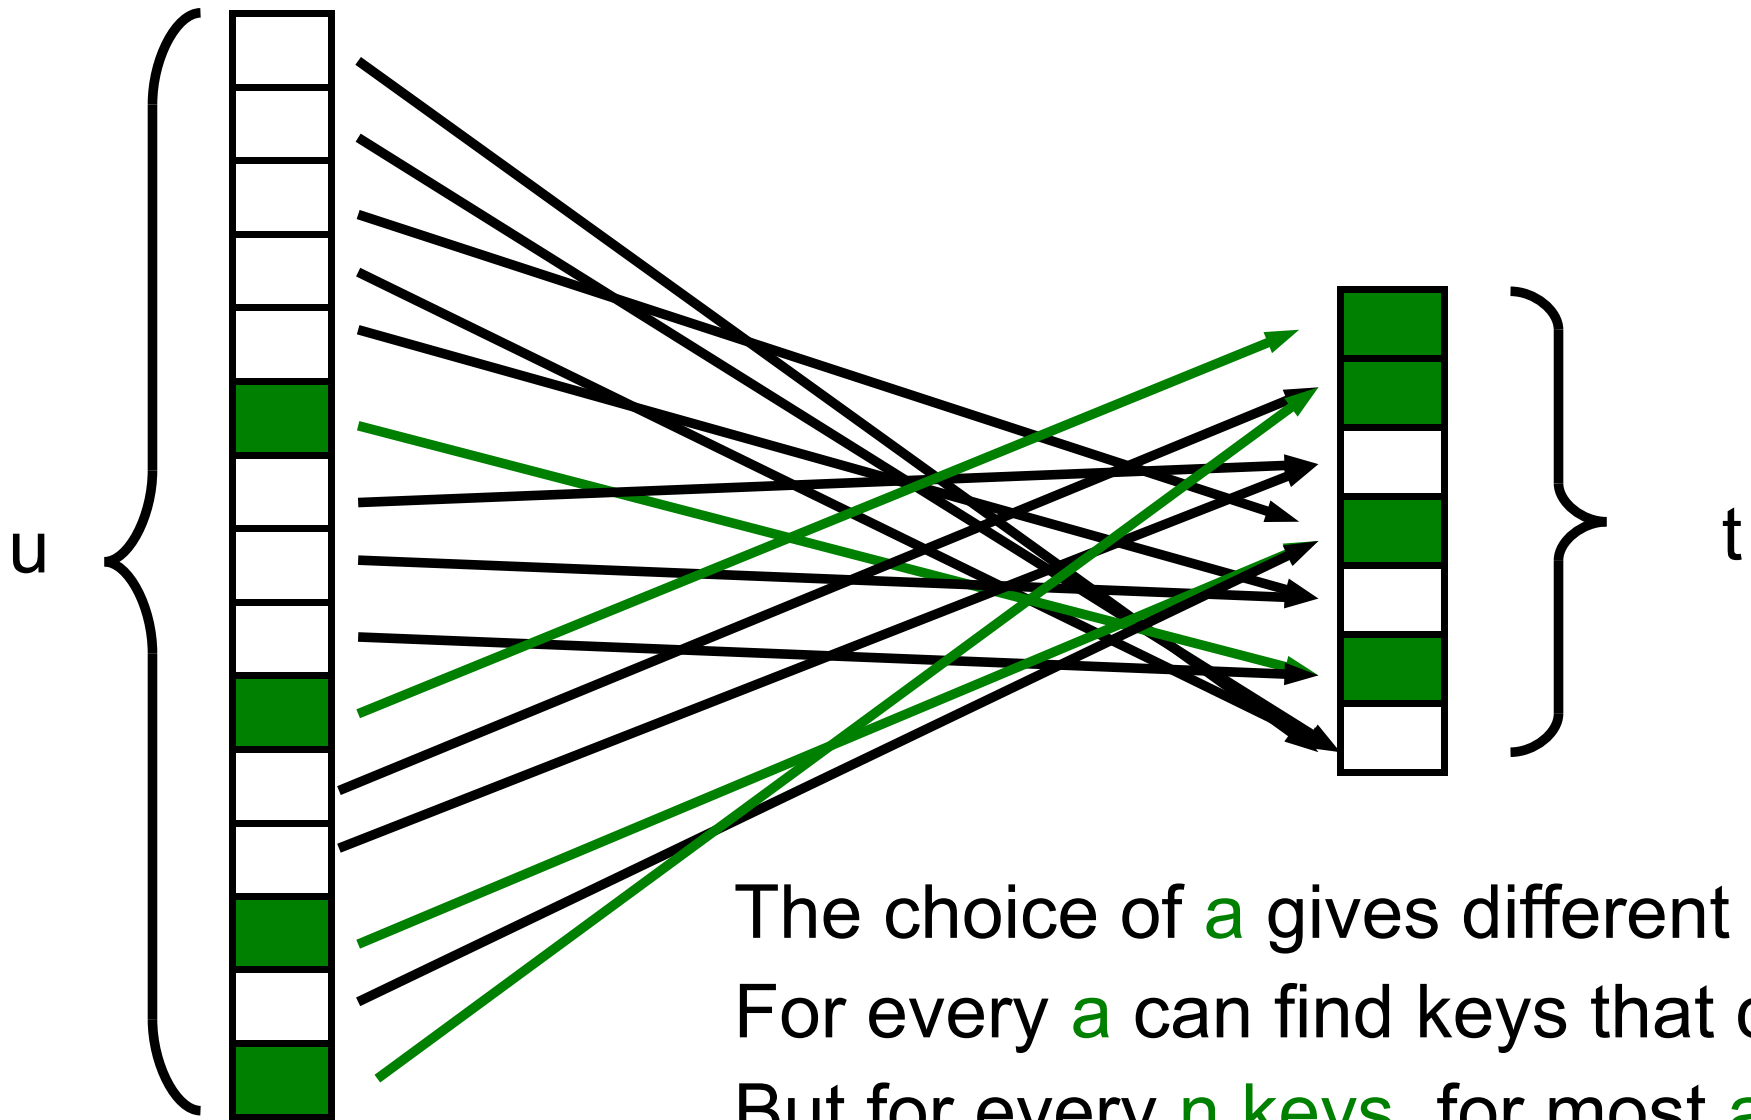


The choice of a gives different arrows
For every a can find **keys that collide**
But for every n keys, for most a
there are no collisions

Keys

Hash function h_a

Table



The choice of a gives different arrows
For every a can find keys that collide
But for every n keys, for most a
there are no collisions

- Want to support INSERT, DELETE, SEARCH for n keys
Keys come from large UNIVERSE = $\{1, 2, \dots, u\}$
We map UNIVERSE into a smaller set $\{1, 2, \dots, t\}$
using a **hash function** $h : \text{UNIVERSE} \rightarrow \{1, 2, \dots, t\}$
- We want that for each of our n keys, the values of h are different, so that we have no collisions
- In this case we can keep an array $S[1..t]$ and
SEARCH(x): ?
INSERT(x): ?
DELETE(x): ?

- Want to support INSERT, DELETE, SEARCH for n keys
Keys come from large UNIVERSE = $\{1, 2, \dots, u\}$
We map UNIVERSE into a smaller set $\{1, 2, \dots, t\}$
using a **hash function** $h : \text{UNIVERSE} \rightarrow \{1, 2, \dots, t\}$
- We want that for each of our n keys, the values of h are different, so that we have no collisions
- In this case we can keep an array $S[1..t]$ and
SEARCH(x): return $S[h(x)]$
INSERT(x): $S[h(x)] \leftarrow 1$
DELETE(x): $S[h(x)] \leftarrow 0$

- Want to support INSERT, DELETE, SEARCH for n keys
Keys come from large UNIVERSE = $\{1, 2, \dots, u\}$
We map UNIVERSE into a smaller set $\{1, 2, \dots, t\}$
using a **hash function** $h : \text{UNIVERSE} \rightarrow \{1, 2, \dots, t\}$
- We want that for each of our n keys, the values of h are different, so that we have no collisions
- Example, think $n = 2^{10}$, $u = 2^{1000}$, $t = 2^{20}$

- Want to support INSERT, DELETE, SEARCH for n keys
Keys come from large UNIVERSE = $\{1, 2, \dots, u\}$
We map UNIVERSE into a smaller set $\{1, 2, \dots, t\}$
using a **hash function** $h : \text{UNIVERSE} \rightarrow \{1, 2, \dots, t\}$
- We want that for each of our n keys, the values of h are **different**, so that we have no collisions
- Can a fixed function h do the job?

- Want to support INSERT, DELETE, SEARCH for n keys

Keys come from large UNIVERSE = $\{1, 2, \dots, u\}$

We map UNIVERSE into a smaller set $\{1, 2, \dots, t\}$

using a **hash function** $h : \text{UNIVERSE} \rightarrow \{1, 2, \dots, t\}$

- We want that for each of our n keys, the values of h are different, so that we have no collisions

- Can a fixed function h do the job?

No, if h is fixed, then one can find two keys $x \neq y$ such that

$h(x)=h(y)$ whenever $u > t$

So our function will use randomness.

Also need compact representation so can actually use it.

- Construction of hash function:

Let t be prime. Write a key x in base t :

$$x = x_1 x_2 \dots x_m \quad \text{for } m = \log_t(u) = \log_2(u)/\log_2(t)$$

Hash function specified by seed element $a = a_1 a_2 \dots a_m$

$$h_a(x) := \sum_{i \leq m} x_i a_i \text{ modulo } t$$

- Example: $t = 97$, $x = 171494$

$$x_1 = 18, x_2 = 21, x_3 = 95$$

$$a_1 = 45, a_2 = 18, a_3 = 7$$

$$\text{Output } 18*45 + 21*18 + 95*7 \text{ mod } 97 = 10$$

- Different constructions of hash function:

Think of hashing s -bit keys to r bits

- Classic solution: for a prime $p > 2^s$, and a in $[p]$,

$$h_a(x) := ((ax) \bmod p) \bmod 2^r$$

Problem: $\bmod p$ is slow, even with Mersenne primes ($p = 2^i - 1$)

- Alternative: let b be a random odd s -bit number and

$$h_b(x) = ((bx) \bmod 2^s) \div 2^{s-r}$$

= bits from $s-r$ to s of integer product bx

Faster in practice. In **C**, think x unsigned integer of $s=64$ bits

$$h_b(x) = (b * x) \gg (s-r)$$

- Analyzing hash functions

The function $h_a(x) := \sum_{i \leq m} x_i a_i \text{ modulo } t$ satisfies

- 2-Hash Claim: $\forall x \neq x', \Pr_a [h_a(x) = h_a(x')] = 1/t$

In other words, on any two fixed inputs, the function behaves like a completely random function

- **n-hash Claim:**

Let h_a be a function from UNIVERSE to $\{1, 2, \dots, t\}$

Suppose h_a satisfies 2-hash claim

If $t \geq 100 n^2$ then for any n keys the probability that two have same hash is at most $1/100$


- **Proof:** $\Pr_a [\exists x \neq y : h_a(x) = h_a(y)]$
 $\leq \sum_{x, y : x \neq y} \Pr_a [h_a(x) = h_a(y)] \quad (\text{union bound})$
 $= \sum_{x, y : x \neq y} \text{?????}$

- **n-hash Claim:**

Let h_a be a function from UNIVERSE to $\{1, 2, \dots, t\}$

Suppose h_a satisfies 2-hash claim

If $t \geq 100 n^2$ then for any n keys the probability that two have same hash is at most $1/100$

- **Proof:** $\Pr_a [\exists x \neq y : h_a(x) = h_a(y)]$
 $\leq \sum_{x, y : x \neq y} \Pr_a [h_a(x) = h_a(y)]$ (union bound)
 $= \sum_{x, y : x \neq y} (1/t)$ (2-hash claim)
 $\leq n^2 (1/t) = 1/100$ 

- So, just make your table size $100n^2$ and you avoid collision
- Can you have no collisions with space $O(n)$?

- Theorem:

Given n keys, can support SEARCH in $O(1)$ time and $O(n)$ space

- Proof:

Two-level hashing:

(1) First hash to $t = O(n)$ elements,

(2) Then hash again using the previous method:

if i -th cell in first level has c_i elements, hash to c_i^2 cells at the second level.

$$\text{Expected total size} \leq E[\sum_{i \leq t} c_i^2]$$

$$= \Theta(\text{expected number of colliding pairs in first level}) =$$

$$= O(n^2 / t)$$

$$= O(n)$$



Queues and heaps

Queue

Operations: ENQUEUE, DEQUEUE

First-in-first-out

Simple, constant-time implementation using arrays:

$A[0..n-1]$

$\text{First} \leftarrow 0$

$\text{Last} \leftarrow 0$

ENQUEUE(x): If $(\text{Last} < n)$, $A[\text{Last}++] \leftarrow x$

DEQUEUE: If $\text{First} < \text{Last}$, return $A[\text{First}++]$

Priority queue

- Want to support
INSERT
EXTRACT-MIN
- Can do it using ??
Time = ?? per query.
Space = ??

Priority queue

- Want to support
INSERT
EXTRACT-MIN
- Can do it using AA trees.
Time = $O(\log n)$ per query.
Space = $O(n)$.
- We now see a data structure that is simpler and somewhat more efficient.
In particular, the space will be n rather than $O(n)$

A binary tree is **complete** if all the nodes have two children except the nodes in the last level.

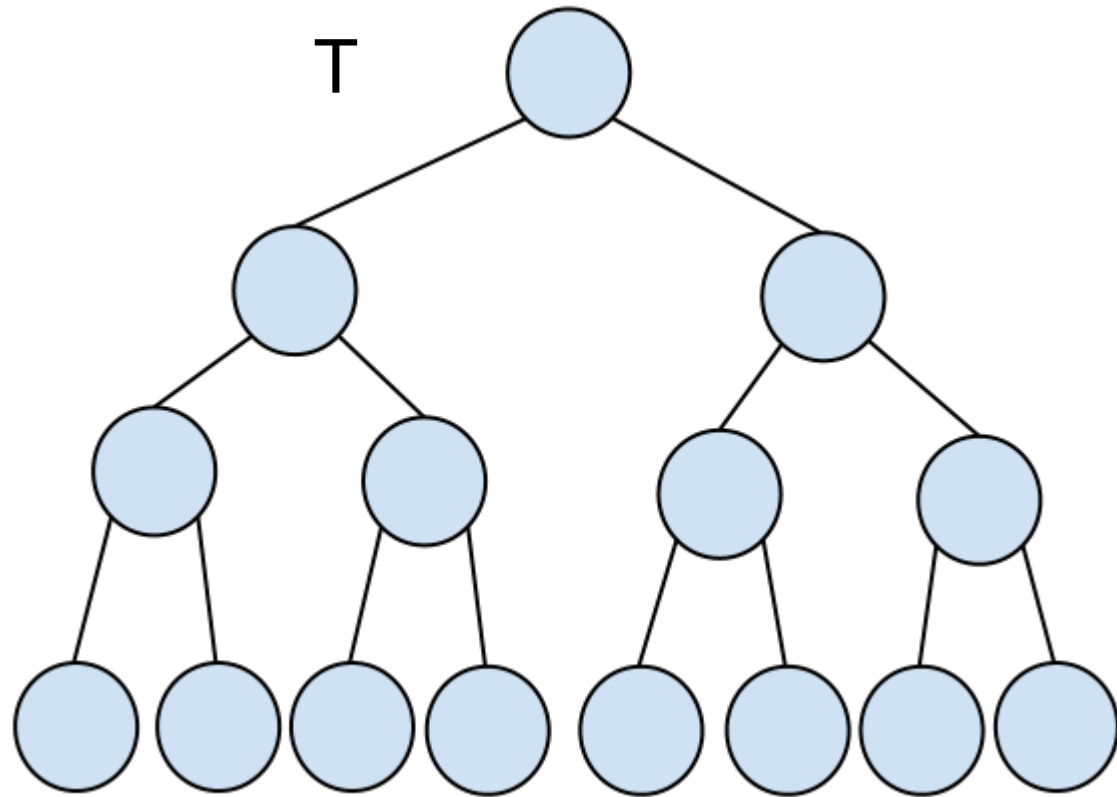
A complete binary tree of depth **d** has **2^d** leaves and **$2^{d+1}-1$** nodes.

Example:

Depth of $T = ?$

Number of leaves in $T = ?$

Number of nodes in $T = ?$



A binary tree is **complete** if all the nodes have two children except the nodes in the last level.

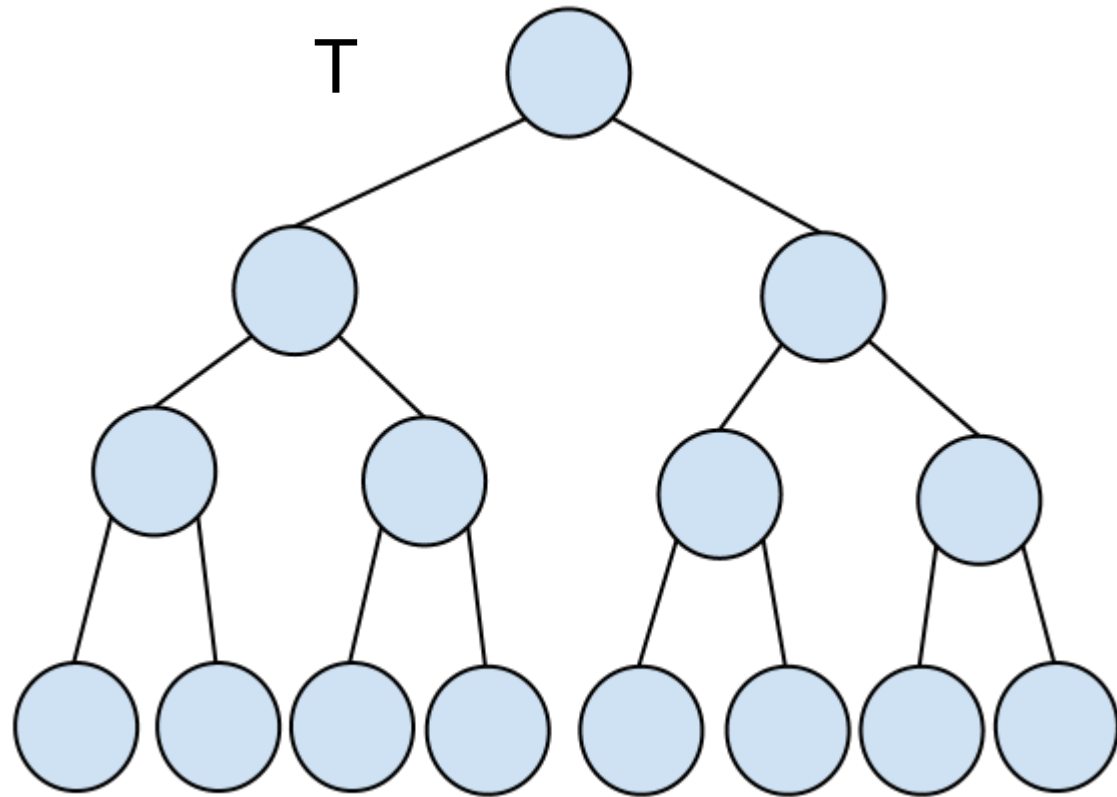
A complete binary tree of depth **d** has **2^d** leaves and **$2^{d+1}-1$** nodes.

Example:

Depth of $T=3$.

Number of leaves in $T=?$

Number of nodes in $T=?$



A binary tree is **complete** if all the nodes have two children except the nodes in the last level.

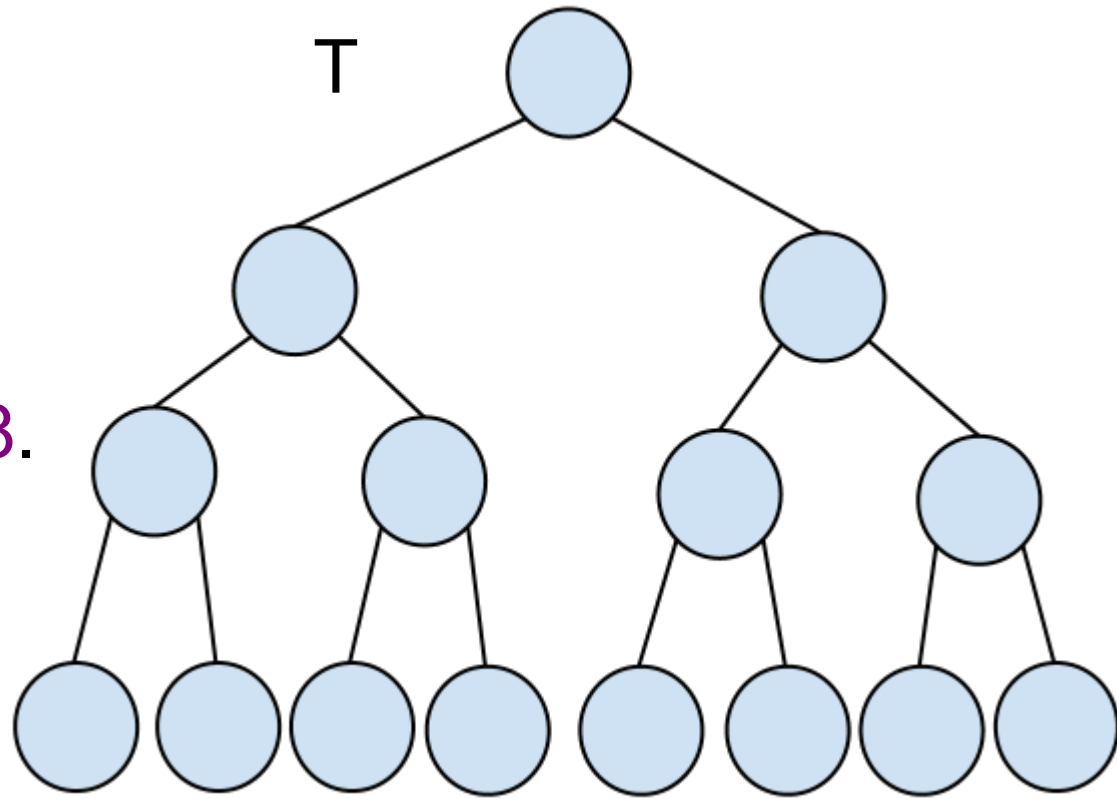
A complete binary tree of depth **d** has **2^d** leaves and **$2^{d+1}-1$** nodes.

Example:

Depth of $T=3$.

Number of leaves in $T=2^3=8$.

Number of nodes in $T=?$



A binary tree is **complete** if all the nodes have two children except the nodes in the last level.

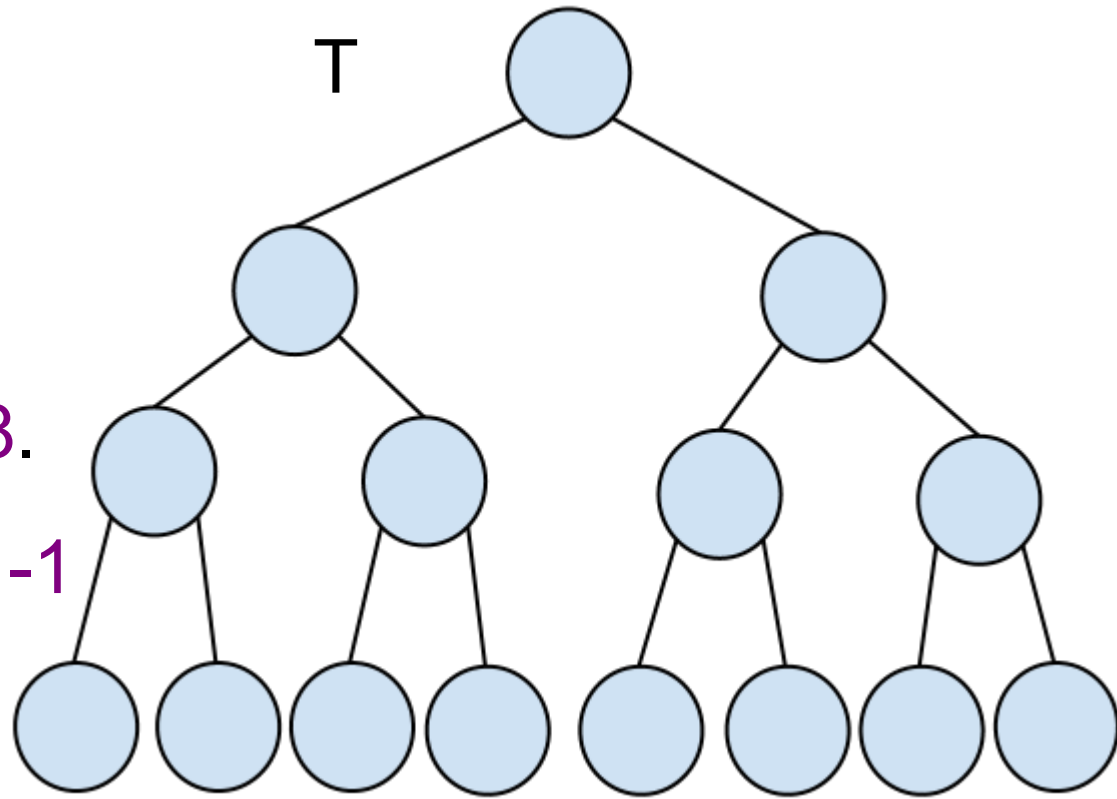
A complete binary tree of depth **d** has **2^d** leaves and **$2^{d+1}-1$** nodes.

Example:

Depth of $T=3$.

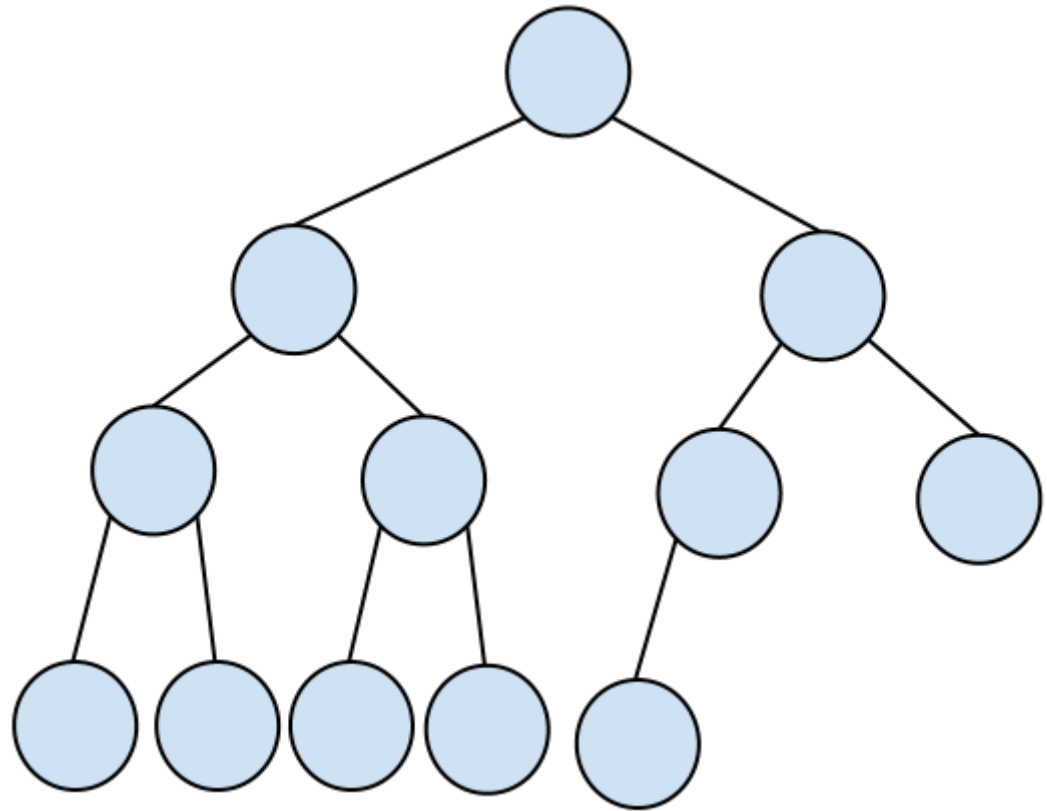
Number of leaves in $T=2^3=8$.

Number of nodes in $T=2^{3+1}-1$
 $=15$.



Heap is like a complete binary tree except that the last level may be missing nodes, and if so is filled from left to right.

Note: A complete binary tree is a special case of a heap.



A heap is conveniently represented using arrays

Navigating a heap:

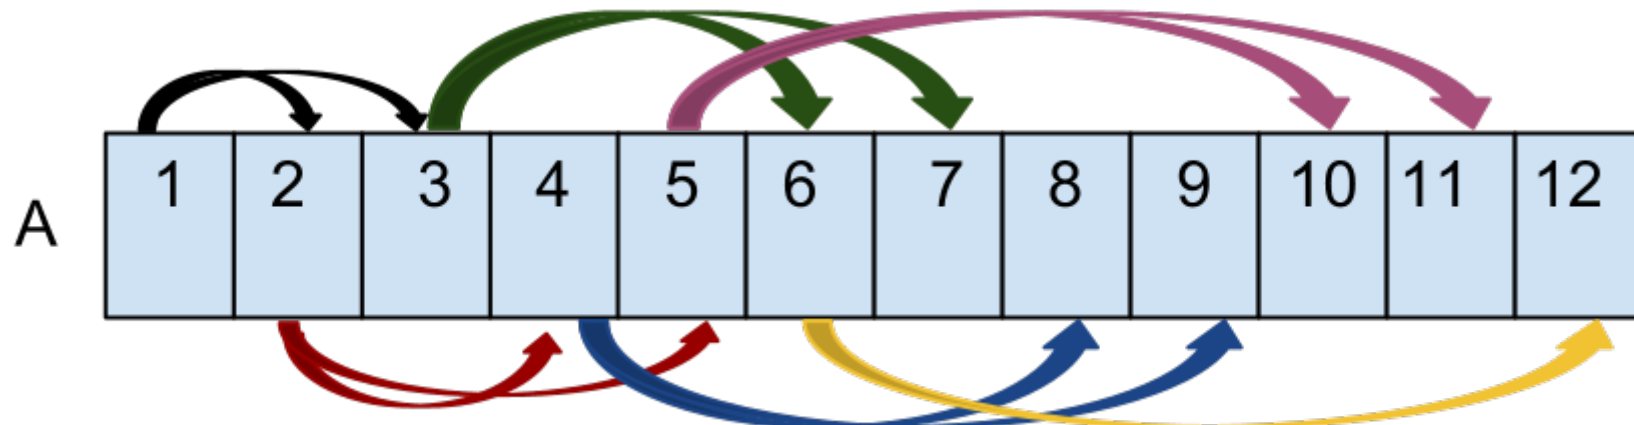
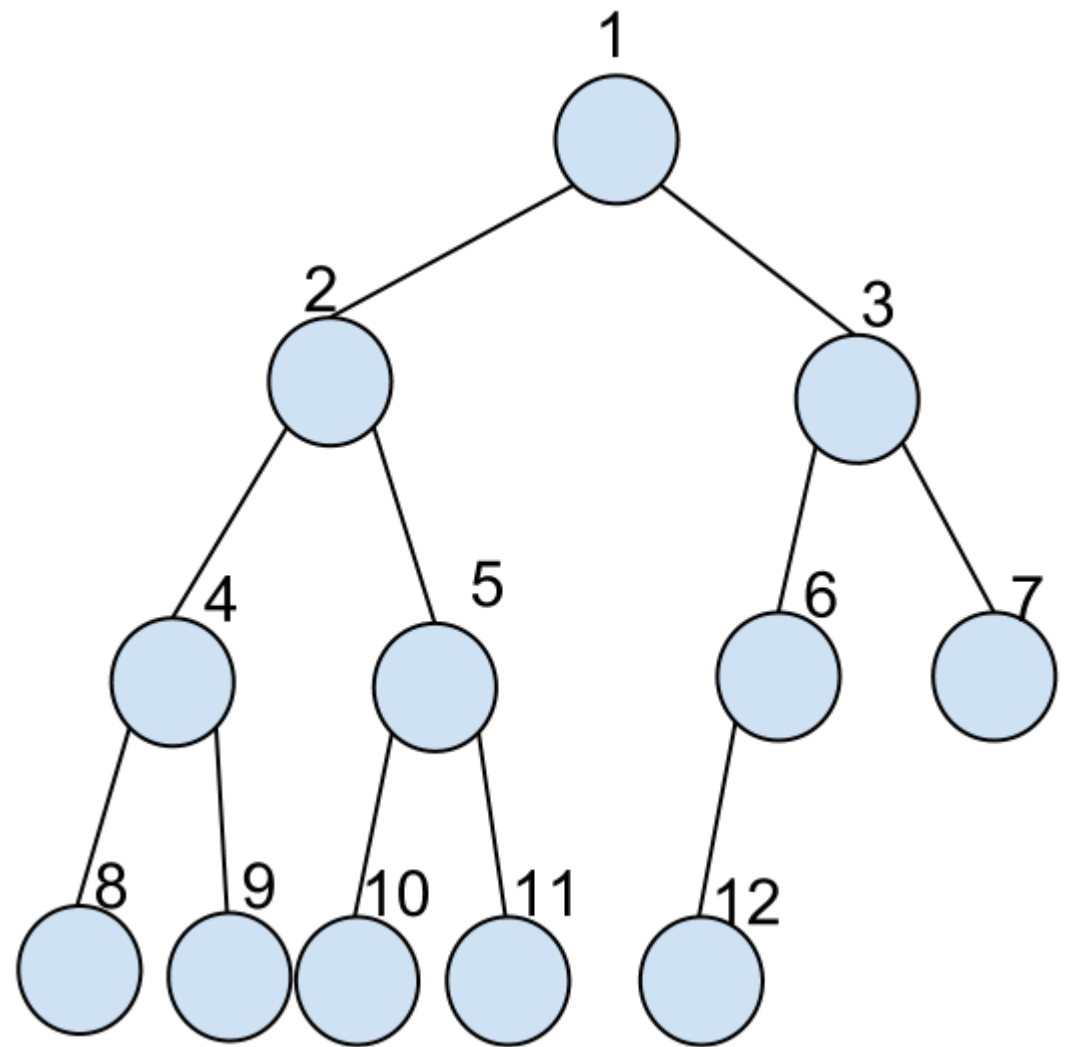
Root is $A[1]$.

Given index i to a node:

$$\text{Parent}(i) = i/2$$

$$\text{Left-Child}(i) = 2i$$

$$\text{Right-Child}(i) = 2i+1$$

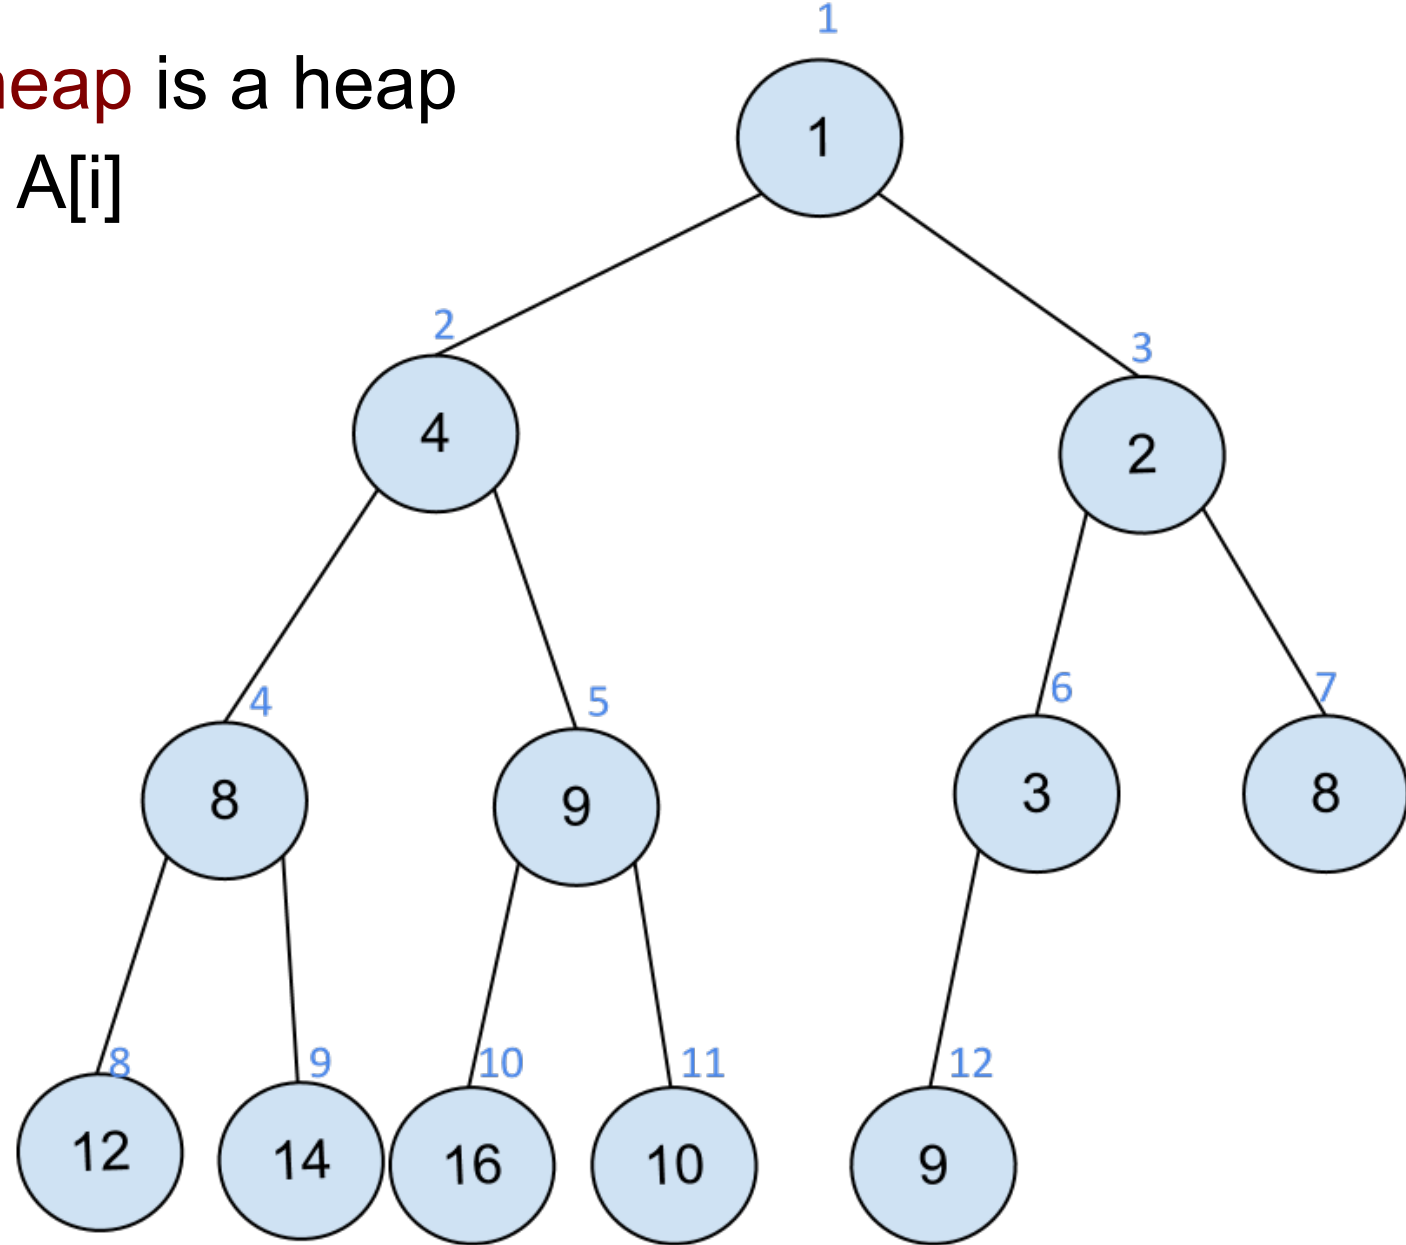


Heaps are useful to dynamically maintain a set of elements while allowing for extraction of minimum (priority queue)

The same results hold for extraction of maximum

We focus on minimum for concreteness.

● **Definition:** A **min-heap** is a heap where $A[\text{Parent}(i)] \leq A[i]$ for every i



	1	2	3	4	5	6	7	8	9	10	11	12
A	1	4	2	8	9	3	8	12	14	16	10	9

Extracting the minimum element

In min-heap A , the minimum element is $A[1]$.

Extract-Min-heap(A)

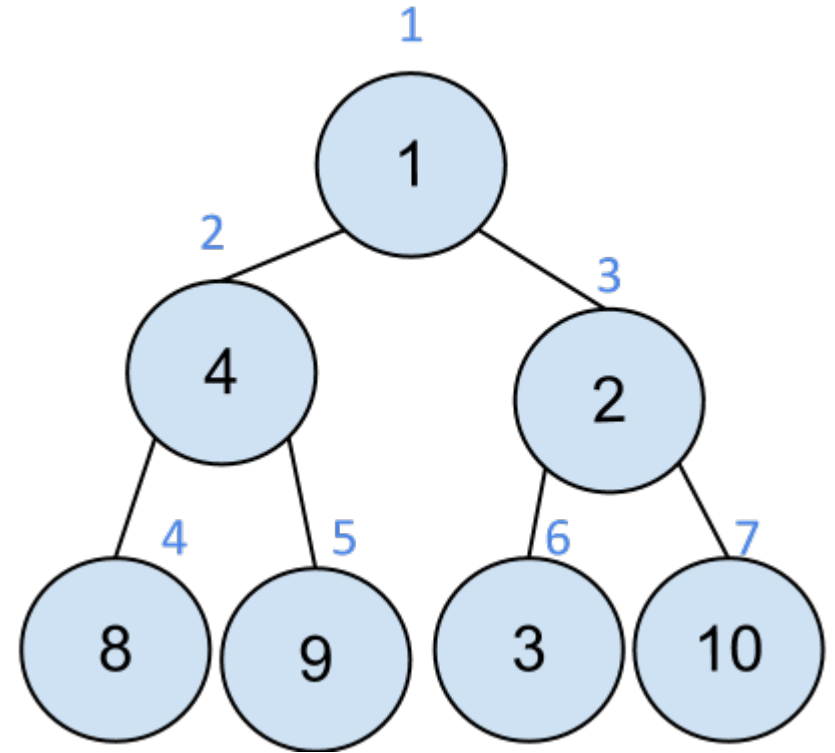
$\text{min} := A[1];$

$A[1] := A[\text{heap-size}];$

$\text{heap-size} := \text{heap-size} - 1;$

Min-heapify($A, 1$)

Return min;



Let's see the steps

Extracting the minimum element

In min-heap A , the minimum element is $A[1]$.

Extract-Min-heap(A)

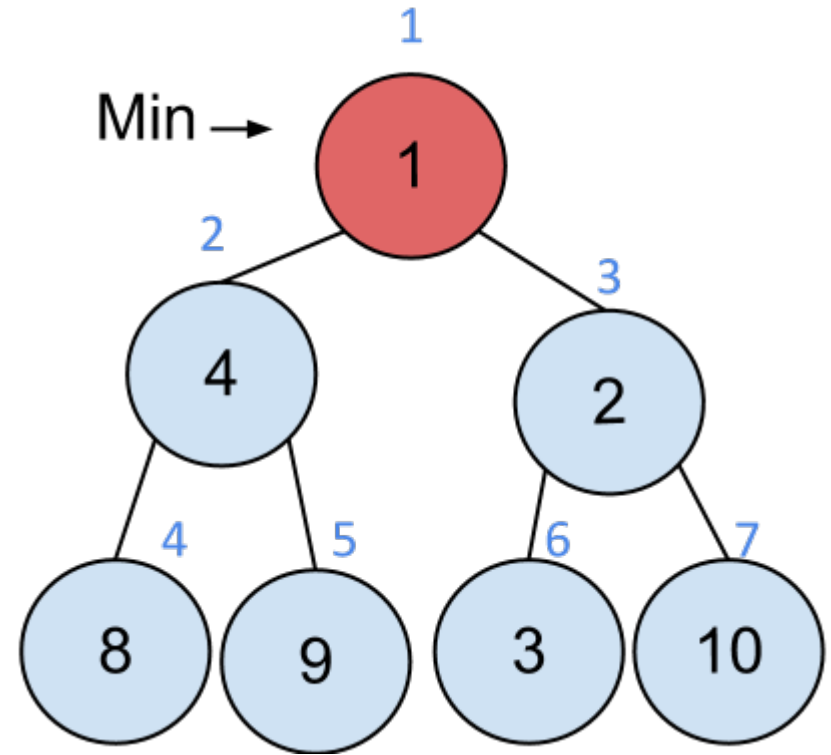
$\text{min} := A[1];$

$A[1] := A[\text{heap-size}];$

$\text{heap-size} := \text{heap-size} - 1;$

Min-heapify($A, 1$)

Return min;



Extracting the minimum element

In min-heap A , the minimum element is $A[1]$.

Extract-Min-heap(A)

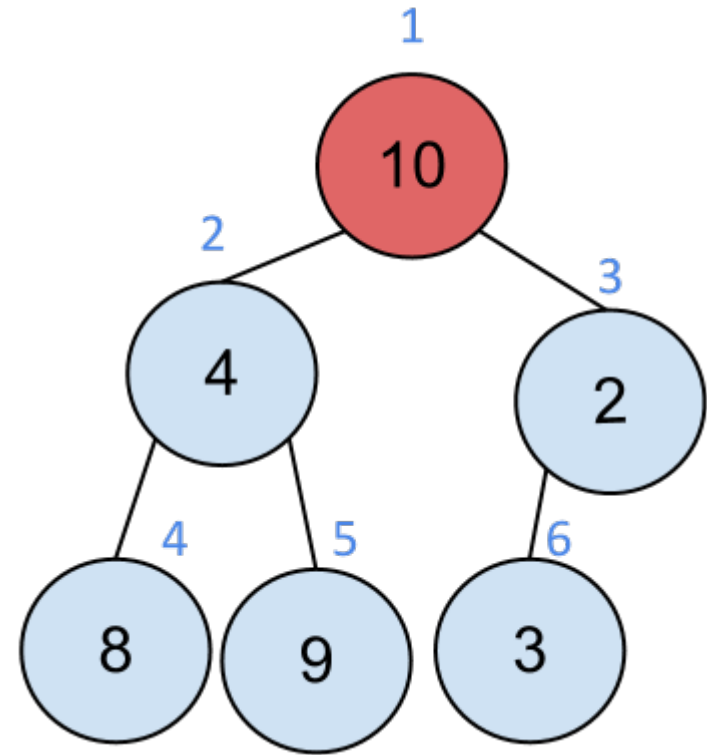
$\text{min} := A[1];$

$A[1] := A[\text{heap-size}];$

$\text{heap-size} := \text{heap-size} - 1;$

Min-heapify($A, 1$)

Return min;



Extracting the minimum element

In min-heap A , the minimum element is $A[1]$.

Extract-Min-heap(A)

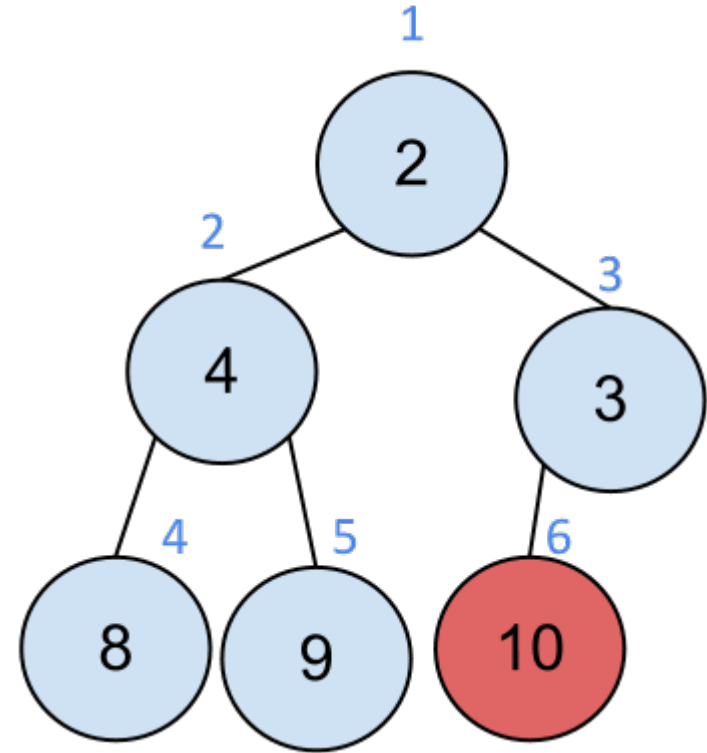
$\text{min} := A[1];$

$A[1] := A[\text{heap-size}];$

$\text{heap-size} := \text{heap-size} - 1;$

$\text{Min-heapify}(A, 1)$

Return min;



Min-heapify is a function that restores the min property

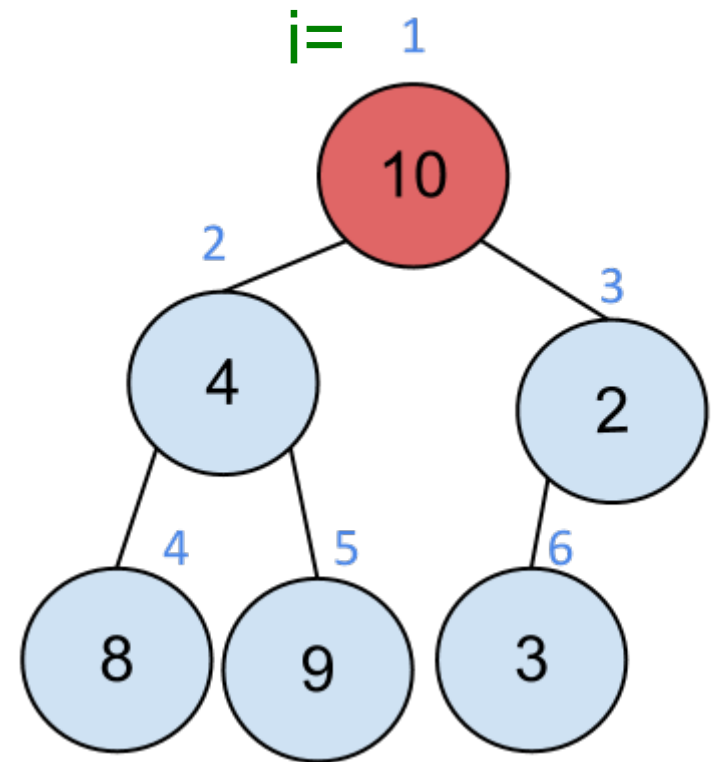
Min-heapify restores the min-heap property

given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

Min-heapify(A, i)

Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {
 exchange $A[i]$ and $A[j]$
 Min-heapify(A, j)
}



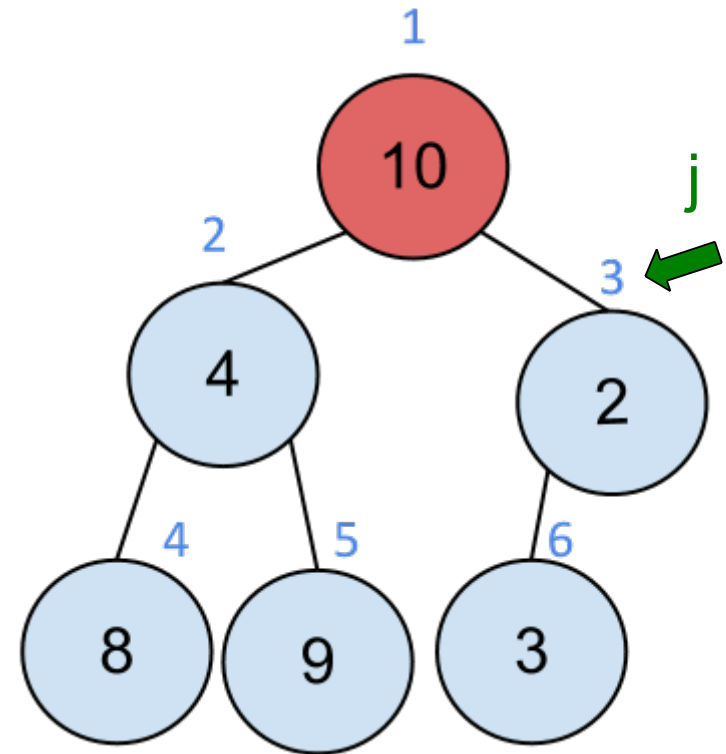
Min-heapify restores the min-heap property

given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

Min-heapify(A, i)

Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {
 exchange $A[i]$ and $A[j]$
 Min-heapify(A, j)
}



Min-heapify restores the min-heap property

given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

Min-heapify(A, i)

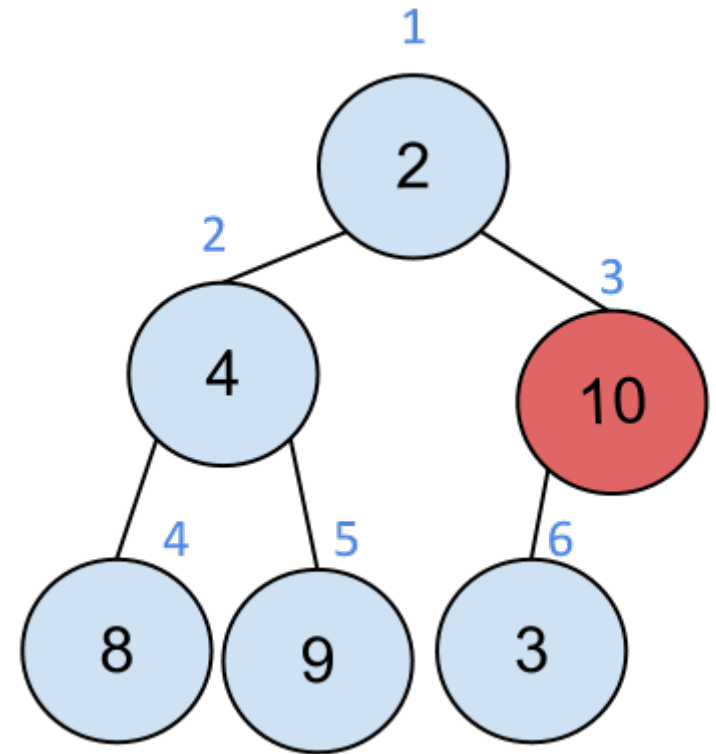
Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {

exchange $A[i]$ and $A[j]$

Min-heapify(A, j)

}



Min-heapify restores the min-heap property

given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

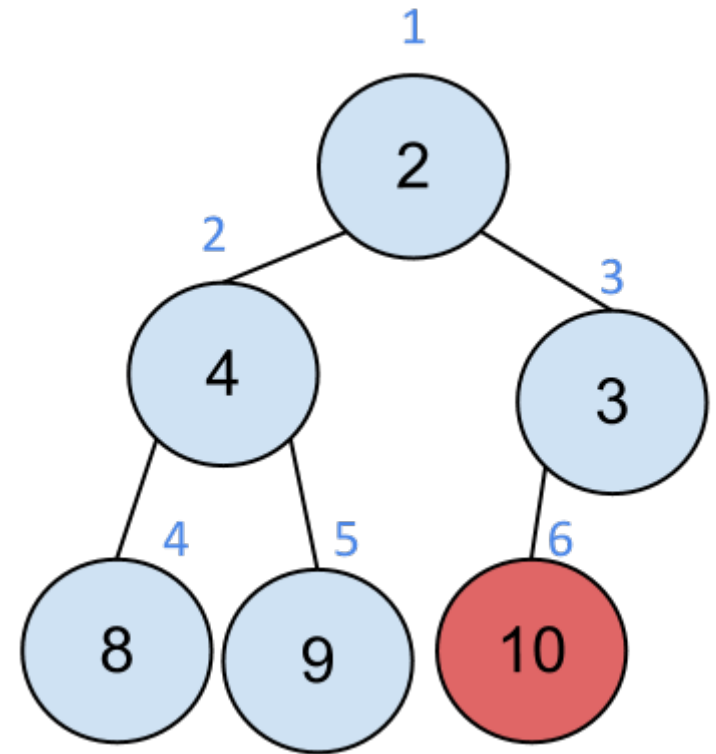
Min-heapify(A, i)

Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {
 exchange $A[i]$ and $A[j]$

 Min-heapify(A, j)

}



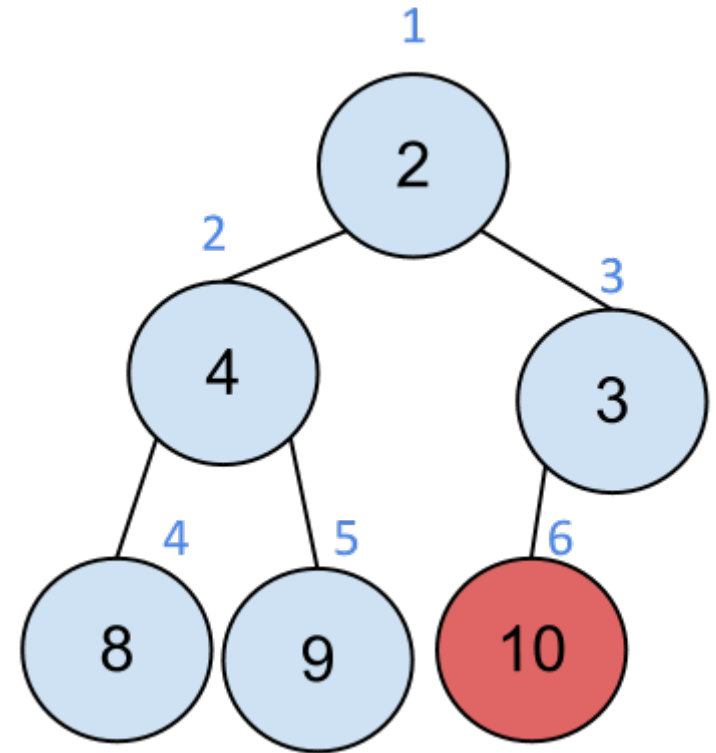
Min-heapify restores the min-heap property

given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

Min-heapify(A, i)

Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {
 exchange $A[i]$ and $A[j]$
 Min-heapify(A, j)
}



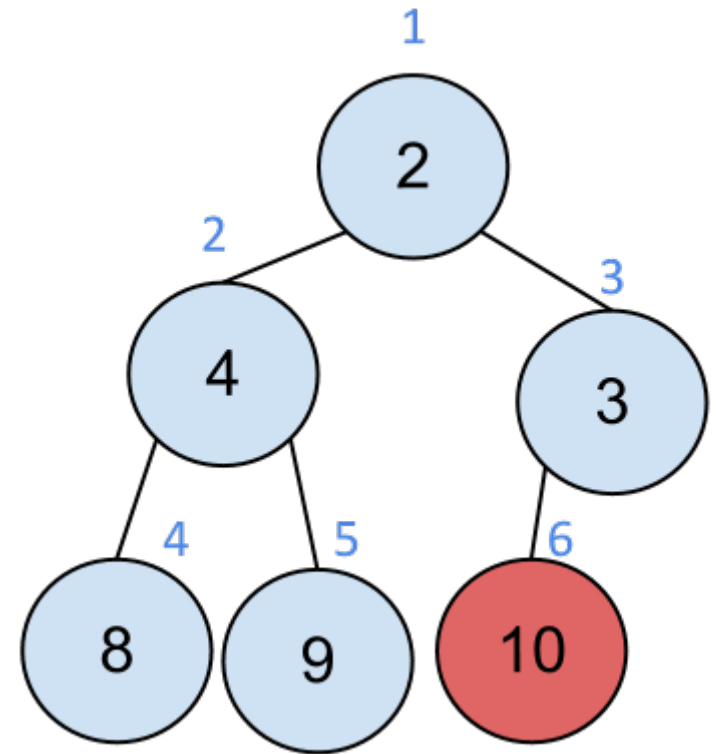
Running time = ?

Min-heapify restores the min-heap property
given array A and index i such that trees rooted at $\text{left}[i]$ and $\text{right}[i]$ are min-heap, but $A[i]$ maybe greater than its children

Min-heapify(A, i)

Let j be the index of smallest node
among $\{A[i], A[\text{Left}[i]], A[\text{Right}[i]]\}$

If $j \neq i$ then {
 exchange $A[i]$ and $A[j]$
 Min-heapify(A, j)
}



Running time = depth = $O(\log n)$

Recall Extract-Min-heap(A)

```
min:= A[1];  
A[1]:= A[heap-size];  
heap-size:= heap-size - 1;  
Min-heapify(A, 1)  
Return min;
```

Hence both Min-heapify and
Extract-Min-Heap take time $O(\log n)$.

Next: How do you insert into a heap?

Insert-Min-heap (A, key)

heap-size[A] := heap-size[A] + 1;

A[heap-size] := key;

for (i := heap-size[A]; i > 1 and A[parent(i)] > A[i]; i := parent(i))
 exchange(A[parent(i)], A[i])

Running time = ?

Insert-Min-heap (A, key)

heap-size[A] := heap-size[A]+1;

A[heap-size] := key;

for($i := \text{heap-size}[A]$; $i > 1$ and $A[\text{parent}(i)] > A[i]$; $i := \text{parent}[i]$)
 exchange($A[\text{parent}(i)]$, $A[i]$)

Running time = $O(\log n)$.

Suppose we start with an empty heap and insert n elements.
By above, running time is $O(n \log n)$.

But actually we can achieve $O(n)$.

Build Min-heap

Input: Array A, output: Min-heap A.

```
For ( i := length[A]/2; i > 0; i -- )  
    Min-heapify(A, i)
```

Running time = ?

Min-heapify takes time $O(h)$ where h is depth.

How many trees of a given depth h do you have?

Build Min-heap

Input: Array A, output: Min-heap A.

```
For ( i := length[A]/2; i > 0; i -- )  
    Min-heapify(A, i)
```

$$\begin{aligned}\text{Running time} &= O\left(\sum_{h < \log n} n/2^h\right) h \\ &= n O\left(\sum_{h < \log n} h/2^h\right) \\ &= ?\end{aligned}$$

Build Min-heap

Input: Array A, output: Min-heap A.

```
For ( i := length[A]/2; i > 0; i -- )  
    Min-heapify(A, i)
```

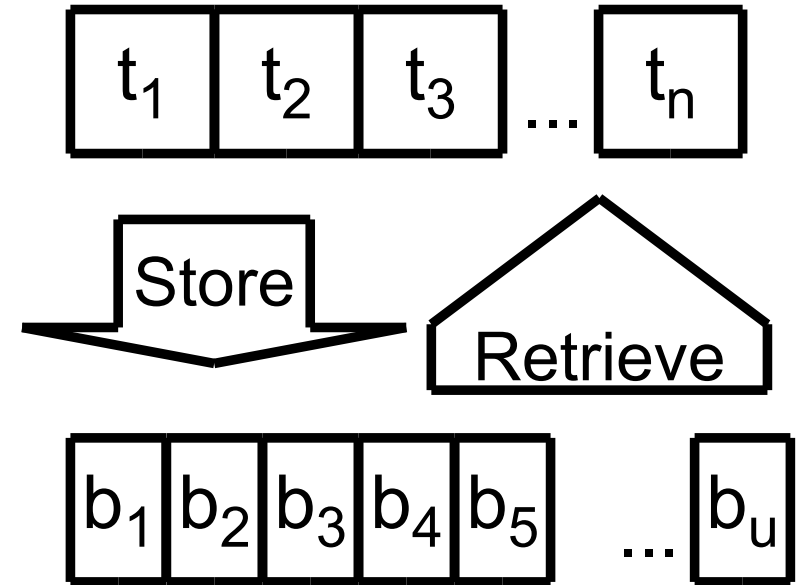
$$\begin{aligned}\text{Running time} &= O\left(\sum_{h < \log n} n/2^h\right) h \\ &= n O\left(\sum_{h < \log n} h/2^h\right) \\ &= O(n)\end{aligned}$$

Next:

Compact (also known as succinct) arrays

Bits vs. trits

- Store n “trits” $t_1, t_2, \dots, t_n \in \{0, 1, 2\}$



In u bits $b_1, b_2, \dots, b_u \in \{0, 1\}$

- Want:

Small space u (optimal = $\lceil n \lg_2 3 \rceil$)

Fast retrieval: Get t_i by probing few bits (optimal = 2)

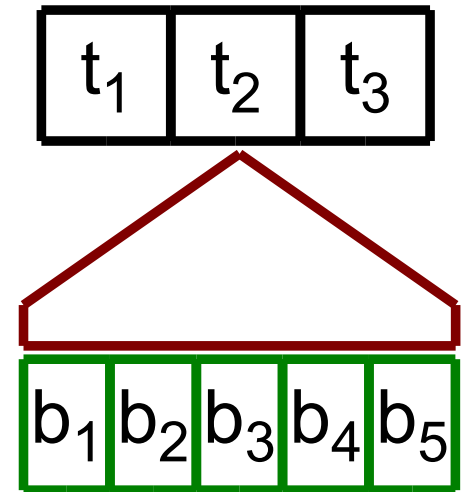
Two solutions

- Arithmetic coding:

Store bits of $(t_1, \dots, t_n) \in \{0, 1, \dots, 3^n - 1\}$

Optimal space: $\lceil n \lg_2 3 \rceil \approx n \cdot 1.584$

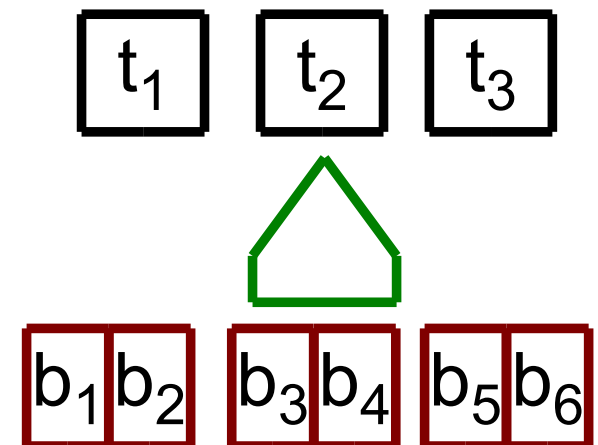
Bad retrieval: To get t_i probe all $> n$ bits



- Two bits per trit

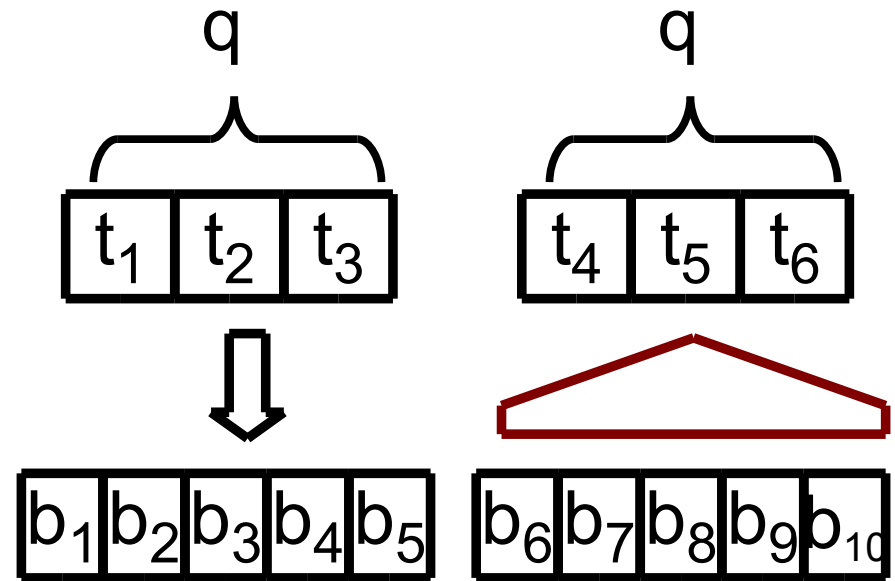
Bad space: $n \cdot 2$

Optimal retrieval: Probe 2 bits



Polynomial tradeoff

- Divide n trits $t_1, \dots, t_n \in \{0,1,2\}$ in blocks of q
- Arithmetic-code each block



$$\begin{aligned} \text{Space: } \lceil q \lg_2 3 \rceil n/q &< (q \lg_2 3 + 1) n/q \\ &= n \lg_2 3 + \textcolor{red}{n/q} \end{aligned}$$

Retrieval: Probe $\textcolor{red}{O(q)}$ bits

**polynomial
tradeoff
between
probes,
redundancy**

Exponential tradeoff

- Breakthrough [Pătraşcu '08, later + Thorup]

Space: $n \lg_2 3 + n/2^{\Omega(q)}$

Retrieval: Probe q bits

exponential
tradeoff
between
redundancy,
probes

- E.g., optimal space $\lceil n \lg_2 3 \rceil$, probe $O(\lg n)$

Delete scenes

Problem: Dynamically support n search/insert elements in $\{0,1\}^u$
Idea: Use function $f : \{0,1\}^u \rightarrow [t]$, resolve collisions by chaining

Function	Search time	Extra space
$f(x) = x$ $t = 2^n$, open addressing	?	?

Problem: Dynamically support n search/insert elements in $\{0,1\}^u$
Idea: Use function $f : \{0,1\}^u \rightarrow [t]$, resolve collisions by chaining

Function	Search time	Extra space
$f(x) = x$ $t = 2^n$, open addressing	$O(1)$	2^u
Any deterministic function	?	?

Problem: Dynamically support n search/insert elements in $\{0,1\}^u$
Idea: Use function $f : \{0,1\}^u \rightarrow [t]$, resolve collisions by chaining

Function	Search time	Extra space
$f(x) = x$ $t = 2^n$, open addressing	$O(1)$	2^u
Any deterministic function	n	0
Random function	? expected	?

Problem: Dynamically support n search/insert elements in $\{0,1\}^u$
Idea: Use function $f : \{0,1\}^u \rightarrow [t]$, resolve collisions by chaining

Function	Search time	Extra space
$f(x) = x$ $t = 2^n$, open addressing	$O(1)$	2^u
Any deterministic function	n	0
Random function	n/t expected $\forall x \neq y, \Pr[f(x)=f(y)] \leq 1/t$	$2^u \log(t)$
Now what? We ``derandomize" random functions		

Problem: Dynamically support n search/insert elements in $\{0,1\}^u$
Idea: Use function $f : \{0,1\}^u \rightarrow [t]$, resolve collisions by chaining

Function	Search time	Extra space
$f(x) = x$ $t = 2^n$, open addressing	$O(1)$	2^u
Any deterministic function	n	0
Random function	n/t expected $\forall x \neq y, \Pr[f(x)=f(y)] \leq 1/t$	$2^u \log(t)$
Pseudorandom function A.k.a. hash function	n/t expected Idea: Just need $\forall x \neq y,$ $\Pr[f(x)=f(y)] \leq 1/t$	$O(u)$

Stack

Operations: Push, Pop

Last-in-first-out

Queue

Operations: Enqueue, Dequeue

First-in-first-out

Simple implementation using arrays.

Each operation supported in $O(1)$ time.